



ENTERPRISE AI WORKFORCE PLATFORM

AI Workforce Commercial & Deployment Playbook

Product · Market · Sales · Pricing · Marketing · AWS Deployment

READ THIS FIRST

Every capability claim in this document is graded against the repository at the commit below.
930 platform tests and 6 browser suites pass. 49 container checks pass on this machine.
Nothing in NOVA has yet run on a real AWS account, and no agent has yet called a real model provider. Those two facts are stated wherever they matter, and they are the top of \$24.

Version 1.0

15 September 2026

Repository state: commit b9319677 (branch main)

Internal master document — not for external distribution as-is

EXECUTIVE SUMMARY

NOVA, honestly, on one page

NOVA is an enterprise AI workforce platform. It deploys governed AI agents into a company's **own** AWS account and connects them to that company's knowledge, systems, communication channels and recurring work — with a policy layer, an approval queue and an audit log that the customer's own auditor can read. It is built on the open-source Hermes agent runtime, which NOVA governs rather than replaces.

PRODUCT

21k lines of NOVA governing a ~600k-line open-source runtime. 13 Control Center screens, 22 channel platforms, 65 MCP servers, 36 grantable plugins, 60 toolsets, FTS5 knowledge retrieval with S3 mirroring, and a Terraform module that builds the whole AWS footprint.

TARGET MARKET

UK accountancy practices with 20–150 staff. 73% of UK firms are turning work away for lack of people (Advancetrack, 2026); MTD for Income Tax went live on 6 April 2026 and turned compliance into daily work; 95% of mid-tier firms expect more AI use inside three years (ICAEW, 2026).

FIRST OFFER

"Client Records Chasing & Response" — one painful, measurable workflow, not "AI agents". Green / Yellow / Red action classes decided before deployment, so the partner knows exactly what the workforce may do unsupervised.

PRICING

Founding customer: £4,500 implementation + £950/month, 12 months, capped scope. Standard SMB from £7,500 + £1,450/month. Priced on the workflow and the governance, never on prompts or tokens. Customer pays AWS directly.

ACQUISITION

100 named UK practices, built by hand from ICAEW/ACCA directories, Companies House and LinkedIn. Multi-touch outreach to the managing partner. 30-day plan with daily targets in \$9. No paid advertising until customer three.

DEPLOYMENT

Single EC2 host in the customer's account, no public inbound, SSM-only access, KMS-encrypted state, IAM with a permissions boundary and per-integration roles assumed with an ExternalId. Full 36-step guide in §17.

The one thing that is not yet true

NOVA has never been deployed to a real AWS account and no agent in it has ever called a real model provider. Everything graded below **LIVE LOCAL PROVEN** was observed running on a developer machine, in containers, against a scripted OpenAI-compatible server. That is a real and useful rung — it caught four defects no unit test could — but it is not a deployment. The first three actions in §24 exist to close exactly this gap, and until they are closed NOVA should be sold as a **paid pilot**, never as a finished product.

CONTENTS

What is in this document

01 How this document was built	6
The evidence ladder	6
What the audit found, in one chart	6
The four validations, kept apart	6
02 What NOVA actually is	8
The architecture, as a customer should see it	8
How the pieces relate	9
Deployment models	9
03 Capability catalogue	10
3.1 AI workforce	10
3.2 Governance	10
3.3 Knowledge (RAG)	11
3.4 Automation	11
3.5 Communication channels	12
3.6 MCP and plugins	12
3.7 AWS deployment	12
3.8 Control Center	13
04 The best first UK market	14
4.1 What the UK market actually looks like	14
4.2 Scoring the candidate verticals	14
4.3 The scoring detail	16
4.4 The decision	16
05 The first NOVA product offer	18
5.1 Why these three, and not the exciting ones	18
5.2 Green / Yellow / Red — decided before deployment, not after	18
5.3 The exact first deployment	19
06 Customer ROI model	20
6.1 The measurement method	20
6.2 The calculation	20
6.3 An illustrative scenario	21
07 Pricing	22
7.1 What the UK market is paying	22

7.2 Recommended pricing	22
7.3 Why the customer pays	23
7.4 What not to do	23
08 Getting the first UK customer	25
8.1 Where the prospects actually are	25
8.2 Prospect criteria	25
09 The first 100 prospects	27
9.1 The tracker	27
9.2 Lead scoring	27
9.3 The 30-day plan	28
10 Marketing and advertising	29
10.1 Organic — where nearly all your effort goes	29
10.2 Direct sales — where the first customer actually comes from	29
10.3 Paid advertising — the recommendation is no	29
11 The NOVA message	31
11.1 Words that are banned, and what to say instead	32
12 Website structure	33
12.1 The homepage	33
12.2 The other nine pages	33
13 The 15-minute sales demo	35
14 The first customer pilot	37
14.1 Success criteria	37
14.2 What kills a pilot	37
15 The UK sales partner model	39
15.1 Recommended structure	39
15.2 When to hire a full-time salesperson	40
16 Commercial best practices	41
17 Deploying NOVA to AWS	43
17.1 What you are actually doing	43
17.2 Every term, explained once	43
17.3 The 36 steps	44
18 AWS cost control	51
18.1 What generates a bill	51
18.2 Price it here, on the day	51

18.3 Beginner checklist for avoiding a surprise bill 52

19 Production checklist 53

19.1 Before any customer 53

19.2 Customer deployment 53

20 12-month commercial roadmap 55

20.1 By function 56

21 Revenue scenarios 57

21.1 The assumptions, stated once 57

21.2 Scenarios 57

21.3 What these numbers hide 58

22 Competitive positioning 59

22.1 Category by category 59

22.2 What is actually differentiated, and how defensible it is 60

23 Risks and honest limitations 61

23.1 Technical and validation risks 61

23.2 Commercial and operational risks 62

24 What I should do next 64

25 Sources and references 66

25.1 UK market and AI adoption 66

25.2 UK accountancy sector 66

25.3 Pricing and commission benchmarks 66

25.4 Regulatory 67

25.5 AWS 67

25.6 Repository — the primary source for every capability claim 68

01 How this document was built

The audit that came before the writing

Nothing in this playbook was written from memory of what NOVA does. The repository was read first — source, tests, Terraform, the Docker build, the Control Center, and eleven internal audit documents — and every capability was placed on a ladder before a single commercial claim was made. That ladder is the repository's own; it is defined in `docs/audits/NOVA_HERMES_CAPABILITY_AUDIT.md` and used throughout the codebase.


The evidence ladder

Read these as increasing strength of evidence. The distinction that matters most commercially is the last one: everything else can be true on a laptop.

Rung	What it means	What it does NOT mean
■ NOT IMPLEMENTED	The capability is absent.	—
■ DECLARED	A config key, schema or type exists.	That anything reads it.
■ WIRED	A runtime call site reads it.	That the runtime acts on it.
■ ENFORCED	The runtime refuses or acts on it, and bypass is not trivial.	That it is covered by a test.
■ TESTED	Covered by an automated test in this repository.	That it has ever run outside a test.
■ PARTIAL	Some of the capability holds; the rest is named and does not.	That the gap is small.
■ LIVE LOCAL PROVEN	Observed running — real processes, real browser, real files — on a developer machine.	That it works on AWS, or with a real model provider.
■ LIVE FIELD PROVEN	Observed running on a real AWS account with a real model provider and real credentials.	Nothing in NOVA has reached this rung yet.

Source: the ladder is the repository's own, from `docs/audits/NOVA_HERMES_CAPABILITY_AUDIT.md` §Ladder.

What the audit found, in one chart



Counted from §3 of this document. The right-hand bar is the finding that shapes every commercial decision here.

The four validations, kept apart

These are four different questions and the document never merges them.

Validation	Status	What was actually done
Architecture	■ TESTED	Boundaries enforced by tests: NOVA depends on stdlib + PyYAML only; runtime imports confined to nova/runtime/<adapter>/. 19 boundary tests.
Implementation	■ TESTED	930 platform tests pass. 6 browser suites (Playwright, real Chromium) pass against a running control plane, each assertion paired against the files on disk rather than against the UI's own state.
Local validation	■ LIVE LOCAL PROVEN	49 container checks: two tenants, two containers, one image, side by side. Auth, RBAC, tenant isolation, CSP, graceful shutdown, no secrets in the image.
AWS validation	■ NOT IMPLEMENTED	No AWS account has been touched. ECR pull, instance profile, KMS volume mount, awslogs delivery, SSM access and the ExternalId integration roles are configured and unexercised.
Real provider validation	■ NOT IMPLEMENTED	No agent has called Bedrock or any hosted model. The live worker run used a scripted OpenAI-compatible server over the runtime's real client.
Production validation	■ NOT IMPLEMENTED	No customer, no production load, no 24/7 soak, no backup restore drill on real infrastructure.

02 What NOVA actually is


In business language, and honestly

A company buys NOVA to put a small team of AI workers inside its own business — workers that can read the company's documents, answer from them, send and receive messages on the company's channels, run work on a schedule, and hand anything consequential to a human before acting. NOVA is the thing that makes that a **governed** arrangement rather than a clever script: who each worker is, what it may touch, what needs a human, and what it did.

Say this to customers, and say it early

NOVA did not write the agent runtime. The execution engine underneath — the task board, the dispatcher, the scheduler, the 22 channel adapters, the plugin system, MCP — is **Hermes**, an open-source project of roughly 600,000 lines. NOVA is about 21,000 lines that sit above it and govern it. This is a strength and should be sold as one: the customer gets a large, actively developed runtime plus a governance and control layer that the runtime does not have and was never trying to be. Claiming to have built all of it is both false and unnecessary.

The architecture, as a customer should see it



Customer systems reached from inside that account:

S3-mirrored knowledge · 22 channel platforms · 65 MCP servers · 36 grantable plugins

NOVA holds no customer credential. Each integration is a separate IAM role assumed with an ExternalId.

Each layer only talks to the one below it. NOVA's own boundary test suite enforces the fourth arrow: only `nova/runtime/<adapter>/` may import the runtime at all.

How the pieces relate

Piece	Who owns it	What it does here
NOVA	You	The tenant bundle (agents, policy, knowledge, channels, automations as declared YAML), the compiler that turns those into runtime configuration, the Control Center, RBAC, the audit log.
Hermes	NousResearch, open source	Executes. Profiles, durable task board, worker processes, cron store, channels, plugins, MCP, tool registry. NOVA never forks it: the patch budget to core is deliberately near zero.
AWS	The customer	Where it all runs. One EC2 host, an encrypted EBS volume, KMS, CloudWatch, IAM, SSM, ECR. Built by NOVA's Terraform module into the customer's account.
Bedrock (or another provider)	The customer	The model. NOVA names the model and the credential variable; the customer supplies the credential. NOVA never stores one.
Customer systems	The customer	Reached through channels (22 platforms), MCP servers (65 in the runtime's curated catalogue) and plugins (36 grantable). Each is a per-agent grant.
Knowledge	The customer	Directories of the customer's own documents, optionally mirrored from an S3 bucket. Indexed into SQLite FTS5 with BM25 ranking — keyword retrieval, not embeddings. Agents cite what they used.
Agents	You configure, customer approves	A NOVA agent <i>is</i> a Hermes profile. That is why per-agent state, credentials, schedules and memory are isolated by directory rather than by a check somebody could forget to write.
Control Center	You	The customer's only surface. 13 screens. The runtime's own web dashboard is an operator console with zero concept of a tenant and must never be exposed to a customer.

Deployment models

Model	Status	What it means
1. Customer-owned AWS (BYOC)	■ TESTED	The Terraform module builds into the customer's account. Their data never leaves it; they hold the KMS key; they pay AWS directly. This is the model to sell. The module is complete and unit-tested; it has not yet been applied.
2. Managed NOVA deployment	■ PARTIAL	The same single-tenant stack, in an account you operate on the customer's behalf. Technically identical; commercially it makes you the data processor and brings a security review, a DPA and an on-call rota you do not yet have.
3. Multi-tenant managed platform	■ NOT IMPLEMENTED	One deployment serving many customers. The code does not support this and the architecture deliberately does not. State is SQLite on one attached volume; the guarantees hold across processes on one host, not across hosts. Tenant isolation today is one tenant per deployment. Do not sell it.

Why single-tenant is the right first answer anyway

The buyer you are targeting is a partner in a professional-services firm who is personally liable for client confidentiality. "It runs in your own AWS account, your data never reaches us, and you can revoke our access by deleting one IAM role" is a far easier sentence to sell than any multi-tenancy story. The architectural limitation and the commercial advantage are the same fact.

03

Capability catalogue

What NOVA can do, graded

Every capability below was verified against source and tests. Where the Control Center shows something, the screen is named. Where a capability exists in the runtime but NOVA does not surface it, that is said rather than implied — a UI placeholder is not a feature.

3.1 AI workforce

Capability	Status	What it does · business value · what production needs
Agents (as runtime profiles)	■ LIVE LOCAL PROVEN	An agent is declared in YAML and compiled into a runtime profile with its own directory, config, persona, policy and credentials. Value: a named worker with a job description the customer can read and change. Production: model credentials per profile.
Agent CRUD from the browser	■ LIVE LOCAL PROVEN	Create, edit, duplicate, archive, restore, delete — all writing to the tenant bundle, then applying. Value: the customer changes the workforce without an engineer. Governance: admin-only, every write audited intent→committed.
Persona / “Soul” editing	■ LIVE LOCAL PROVEN	The standing instruction each agent carries on every turn, edited through the real backend and capped at 20,000 characters. Written to the bundle, never to the derived SOUL.md — which apply overwrites.
Objectives → task graph	■ TESTED	A declared objective compiles into ordered steps on the runtime's board. Value: multi-step work a human can watch. Gap: the full dependency graph is in the runtime; NOVA surfaces steps, not the graph.
Durable task board	■ LIVE LOCAL PROVEN	Tasks and run history survive process and container restarts (runtime's SQLite Kanban store). Surfaced on the Work screen, tenant-scoped.
Task detail (runs, comments, artifacts)	■ TESTED	Per-task history. Security: the artifact's host path is deliberately never returned by the API.
Worker processes / isolation	■ LIVE LOCAL PROVEN	Each agent runs in its own OS process with its own working directory. Proven by running real dispatched workers, which is how four seam defects were found.
Supervisor / routing	■ TESTED	Objectives route to agents by declared capability; delegation is validated against the roster at load, so a typo fails at configuration time.
Agent memory	■ WIRED	The runtime keeps curated per-profile memory files. NOVA surfaces nothing. Do not sell agent memory.
Agent status	■ LIVE LOCAL PROVEN	Enabled/disabled, applied-vs-saved, profile materialised, scheduler liveness. Only states backed by real runtime data are shown.

3.2 Governance

Capability	Status	What it does · business value · what production needs
RBAC (viewer / admin)	■ LIVE LOCAL PROVEN	Two roles; an exact-match route table where anything undeclared requires admin. Value: staff can watch, partners can change. Proven in-container: viewer 403s on policy, decisions, budget and every write.
Named principals, hashed tokens	■ LIVE LOCAL PROVEN	nova token new mints a token, prints it once and stores only its SHA-256. NOVA never holds the credential. The audit log records the human, not a service name.
Policy compile → decide	■ LIVE LOCAL PROVEN	A tenant policy compiles per agent into a document the enforcement plugin reads at every tool call. Fail-closed: a missing or unreadable policy denies, it does not fall open.
Tool permissions & deny rules	■ ENFORCED	Denials compile into the runtime's own unconditional deny list, evaluated before any bypass mode. Honest limit: positive scoping (toolsets/allow) is recorded but not enforced by the adapter, and the materialiser says so.
Approvals queue	■ TESTED	Actions a deployment requires a human to approve stop and queue. Release and refusal are admin-only and audited.

Capability	Status	What it does · business value · what production needs
Audit logging	■ LIVE LOCAL PROVEN	Write-ahead: <i>intent</i> before the act, <i>committed</i> or <i>failed</i> after, with the authenticated human and a correlation id. Rotation, retention and tamper-evidence implemented. Value: this is what a partner's PI insurer asks for.
Tenant isolation	■ LIVE LOCAL PROVEN	One tenant per deployment; per-agent state isolated by directory. Proven with two containers side by side: each token 401s against the other tenant, automations and provenance are invisible across the boundary.
Automation governance	■ TESTED	An automation is a <i>spec</i> that compiles — its permissions, knowledge and channels must be a subset of the owning agent's. There is no path from the Control Center to a raw scheduled prompt.
Per-channel approval policy	■ TESTED	A channel can require approval for actions that would otherwise pass — e.g. anything that sends externally.
Capability grants (MCP / plugins)	■ LIVE LOCAL PROVEN	Per-agent grants declared in the bundle and compiled by apply, so they survive re-application. Admin-only, audited.

3.3 Knowledge (RAG)

Capability	Status	What it does · business value · what production needs
Corpora & ingestion	■ LIVE LOCAL PROVEN	A corpus is a directory plus include/exclude globs, a size cap and a classification label. The walk refuses symlinks, anything resolving outside the root, and dot-directories — this is the security boundary of the capability.
Retrieval architecture	■ TESTED	SQLite FTS5 , porter stemmer, BM25 ranking. No embeddings, deliberately. Sell it as: keyword search that cites its source. Do not sell: semantic/vector search. Say so before a technical buyer asks.
Browser upload	■ LIVE LOCAL PROVEN	Documents uploaded through the Control Center, name rebuilt not trusted (traversal and Unicode-mangling both tested), re-indexed on upload so “stored” never silently means “invisible”.
S3-mirrored corpora	■ TESTED	A corpus declares an S3 origin; NOVA mirrors it and the existing ingest runs over the result. Keys escaping the root, unaccepted types and oversized objects are refused and named. Proven against a real boto3 client over an in-process S3 server — not on real S3.
Corpus permissions	■ ENFORCED	An agent reads only corpora its spec grants, validated at bundle load. Declaring a corpus also grants the search tool — a trap found by running a live worker.
Citations	■ TESTED	Retrieved chunks carry their document, so an answer can name where it came from. Value: in a professional-services firm an uncited answer is unusable.
Document types	■ PARTIAL	Text and Markdown are first-class. Extraction of richer formats is a runtime capability NOVA reports via a capability flag; it has not been exercised end to end. Scope PDFs explicitly in a pilot, do not assume them.

3.4 Automation

Capability	Status	What it does · business value · what production needs
Governed automations	■ TESTED	Declared, compiled, then created in the runtime's own cron store under the owning agent's profile. The objective is capped at 4,000 characters.
Pause / resume / delete	■ LIVE LOCAL PROVEN	Admin-only, audited, surfaced per agent and on the Automations screen.
Automation provenance	■ LIVE LOCAL PROVEN	A tenant-salted digest registry records what NOVA created and why. Survives a container restart — verified.
Execution ledger	■ WIRED	The runtime keeps per-execution rows. NOVA reads the agent's own ledger directly (the runtime's own accessor resolves to the wrong store — a real bug found and worked around). Surfaced against seeded data only.
Schedules actually firing	■ NOT IMPLEMENTED	Read this twice. The cron ticker lives inside the runtime's gateway; there is no standalone daemon. A deployment can hold a perfectly correct schedule that nothing executes . The Automations screen leads with scheduler liveness and says “nothing is running these schedules” when nothing is. Running a gateway process is a deployment step that has never been done, and it blocks any pilot that promises scheduled work.

3.5 Communication channels

Capability	Status	What it does · business value · what production needs
Provider catalogue	■ LIVE LOCAL PROVEN	22 platforms, read from the runtime's own plugin manifests rather than a hand-kept list. Each carries how its support was established.
Verification status	■ DECLARED	Zero of the 22 are field-validated. 5 were read in source (Slack, Discord, Telegram, Email, WhatsApp); 17 are manifest-only. A test asserts that no provider may claim field validation, so promoting one requires evidence. Never put a channel logo on a slide without connecting it first.
Grant enforced twice	■ ENFORCED	NOVA's parser refuses a route outside the allowed agents, and the grant compiles into the runtime's own fail-closed allowlist. A validation you can bypass by editing YAML is advice; a runtime that will not deliver is a control.
Credential isolation	■ PARTIAL	Stated precisely, because the nuance matters: credentials in a per-agent store are isolated. A credential exported into the <i>host</i> environment is readable by every agent on that host — verified. Isolation is a deployment property, and the readiness report says where each credential resolved from.
No secret can enter a declaration	■ ENFORCED	token, api_key, secret, password, credential and app_secret are refused at parse time at any nesting depth, and stripped from what NOVA writes.
Webhook transport	■ DECLARED	WhatsApp and similar need a publicly reachable HTTPS endpoint. The AWS module has no inbound rules at all . That is a deliberate deployment decision nobody has made yet.

3.6 MCP and plugins

Capability	Status	What it does · business value · what production needs
MCP catalogue	■ LIVE LOCAL PROVEN	65 curated servers from the runtime's pinned manifests. Granting is per agent, declared in the bundle, compiled by apply. Tools arrive as a toolset named <code>mcp-<server></code> and are decided by the same policy as any other tool.
OAuth reality	■ PARTIAL	54 of the 65 authenticate with OAuth , which needs a browser consent the Control Center cannot perform. A granted OAuth server is reported as <i>configured, not yet authorized</i> with the exact command that completes it. 10 need no auth and work immediately.
Catalogue-only security model	■ ENFORCED	Arbitrary MCP servers cannot be added. A stdio server is a command line; accepting one from a browser form would be remote code execution with a form around it. The one catalogue entry that clones and builds is refused too.
Plugin grants	■ LIVE LOCAL PROVEN	36 grantable of 58 discovered. Three states, because the runtime has three: bundled backends load unless disabled, so “not listed” and “disabled” are different answers. Channels are excluded — they have their own screen.
Plugin enable merge safety	■ TESTED	A plugin grant merges into the enabled list rather than replacing it. Replacing it would silently drop NOVA's own policy plugin — the most dangerous state a governance control can be in, because it passes review by inspection. This actually happened once and is now pinned by a test.

3.7 AWS deployment

Capability	Status	What it does · business value · what production needs
Terraform module	■ TESTED	13 resources: EC2, EBS, KMS key + alias, CloudWatch log group, security group (egress 443 only, no ingress), 2 IAM roles, 2 role policies, a permissions boundary policy, an attachment and an instance profile. 16 input variables.
Docker image	■ LIVE LOCAL PROVEN	304 MB, non-root uid 10001, no apt layer, health check in the standard library, immutable tagging, state on a declared volume. 49 checks pass including 8 for “no secrets in the image”.
ECR	■ DECLARED	Not created by the module; the instance role is granted pull. Nothing has been pushed to any registry.
IAM design	■ TESTED	A runtime role with no customer-data permissions , under an explicit permissions boundary so a future mistaken attachment cannot widen it. One role per integration, each assumed with <code>sts:ExternalId</code> .
SSM Session Manager	■ DECLARED	The only way in. No inbound port, no SSH key, no bastion. Configured, never used.
IMDSv2 required	■ TESTED	Asserted by the module's tests. Closes the classic SSRF-to-credentials path.

Capability	Status	What it does · business value · what production needs
KMS encryption	■ TESTED	Volume and logs. The customer may supply their own key ARN — and for a professional-services buyer, they should.
Bedrock	■ DECLARED	Model ids enumerated, never wildcarded. No model has ever been invoked.
Secrets Manager	■ DECLARED	Readable under one prefix and nowhere else.

3.8 Control Center

13 screens. Everything a customer sees comes through NOVA's own tenant-scoped control API — never the runtime's operator dashboard, which has no concept of a tenant at all.

Capability	Status	What it does · business value · what production needs
Overview · Agents · Work · Objectives	■ LIVE LOCAL PROVEN	The daily surface. Agent profile pages carry Overview, Work, Soul, Model, Capabilities, Schedules, Activity, Knowledge, Channels, Extensions, Permissions and Usage.
Approvals · Activity	■ TESTED	The queue, and every refusal and escalation the policy layer recorded. Permitted calls are deliberately not logged as decisions.
Knowledge · Channels · Automations	■ LIVE LOCAL PROVEN	Corpora with upload, sync and reindex; the channel catalogue with per-agent grants; schedules with liveness stated first.
Policies · Usage	■ TESTED	Admin-only. Usage reports observed token and estimated cost.
Usage is not a spending cap	■ PARTIAL	Never sell it as one. Run-budget seconds trigger a soft wrap-up message; nothing terminates. Token and cost figures are observed, not enforced. A hard spending ceiling is structurally unavailable and the repository says so.
Settings & branding	■ LIVE LOCAL PROVEN	Company details, brand colour, logo — applied across the dashboard and into every agent's persona, so agents introduce themselves as the customer's product. White-labelling is real and works.
Security posture	■ LIVE LOCAL PROVEN	Strict CSP (no inline script, no external anything), X-Frame-Options DENY, bearer auth, cross-origin writes refused, form-encoded writes refused.

04

The best first UK market

Researched, scored, chosen

The choice below is made on evidence about the UK market in 2026, not on which sector sounds most exciting. Sources are cited inline and listed in §25. Where a figure is an estimate rather than a published statistic, it says so.

4.1 What the UK market actually looks like

Finding	Figure	Source	Type
UK businesses (10+ employees) using at least one AI technology	~35%, up from ~12% in Sept 2023	ONS, <i>Artificial intelligence in UK businesses</i> , 20 July 2026	Published
All size bands	29%	ONS, same release	Published
Adoption is shallow : average AI technologies per adopting business	~1.4 → ~1.6 since late 2023	ONS, same release	Published
UK SMEs using AI in some form / using it extensively to automate operations	54% / 11%	British Chambers of Commerce, March 2026	Published
AI-using businesses reporting increased revenue	12%	ONS / DSIT, 2026	Published
UK SMEs (0–249 employees)	~5.68 million	UK Business Population Estimates 2025, via House of Commons Library	Published
UK accounting & auditing businesses	~29,600 (IBISWorld) to ~35,900 (Companies House–derived)	IBISWorld 2025; Firmsbase 2026	Published, two definitions
UK accounting firms turning away clients for lack of staff	73%	Advancetrack <i>Accounting Talent Index</i> 2026	Published
Firms saying sustained workloads could push people out of the profession	74%	Advancetrack, same index	Published
Mid-tier firms expecting increased AI use within 3 years	95%	ICAEW mid-tier firms research, 35 firms, Feb–Mar 2026	Published
Firms whose strategy explicitly includes AI adoption	86%	ICAEW, same research	Published
Firms that believe they can assess AI's workforce impact	17%	ICAEW, same research	Published
MTD for Income Tax live for sole traders/landlords above £50,000	from 6 April 2026	HMRC Making Tax Digital timetable	Published

The single most useful number in this table


54% of UK SMEs use AI in some form; 11% use it extensively to automate operations (BCC, March 2026), and adoption is shallow — about 1.6 technologies per adopting business (ONS, July 2026). The market has bought chat assistants and stopped. Almost nobody has deployed a governed workforce that touches real systems. That is precisely the gap NOVA is built for, and it means the sales conversation is *not* “have you considered AI” — it is “you tried AI, it saved drafting time, and the admin came straight back. Here is why.”

4.2 Scoring the candidate verticals

Twelve factors, each scored 1–10, weighted by how much they decide a *first* customer rather than a tenth. Weights are my judgement and are stated so you can disagree with them; the underlying observations are sourced above and in

§25.

Factor	Weight	Why it is weighted that way
Repetitive, documentable workflows	1.5×	Decides whether NOVA can do the job at all.
ROI measurability	1.5×	You need a case study more than you need revenue.
Accessible decision maker	1.4×	Owner/partner who can sign without a committee.
Willingness to pay	1.3×	Firm must already buy software and outsourced labour.
Document volume	1.2×	Feeds the knowledge capability, which is real today.
Communication volume	1.2×	Where the hours actually go.
Implementation time	1.1×	A pilot that takes four months kills momentum.
Case-study likelihood	1.1×	Will they let you name them?
AI adoption / readiness	1.0×	Warm enough to answer, not so warm it is crowded.
Technical readiness	0.9×	Existing CRM/practice software to integrate with.
Integration complexity (inverted)	0.9×	Fewer bespoke connectors to build.
Regulatory risk (inverted)	1.2×	A regulated buyer adds months of review.



Weighted composite of the twelve factors above. Scores are my assessment against the evidence in §4.1 and §25 — they are a decision tool, not a measurement.

4.3 The scoring detail

Vertical	Wkflw	ROI	DM	Pay	Docs	Comm s	Impl	Case	Adopt	Tech	Integ	Reg	Score
Accountancy practices	9	9	9	8	9	9	8	9	8	8	7	6	8.4
Property / block management	9	8	8	7	7	9	7	7	6	6	7	7	7.5
Recruitment agencies	9	8	8	7	6	9	8	7	7	7	7	6	7.3
IT managed service providers	8	8	8	7	6	8	7	6	8	9	8	6	7.1
Digital / marketing agencies	8	7	9	6	6	8	8	7	9	8	7	5	6.8
Logistics & freight	8	8	6	7	7	8	6	6	5	6	6	6	6.5
Professional services (other)	7	7	7	7	8	7	7	6	6	6	6	6	6.4
E-commerce operations	8	8	7	5	4	9	7	6	8	8	6	5	6.0
Legal firms	8	6	5	8	9	7	4	5	5	5	5	3	5.2
Financial services (FCA)	8	6	3	9	9	7	3	4	6	7	4	1	4.4

Columns in the order of the weight table in §4.2. Regulatory and integration columns are already inverted (10 = low risk / low complexity).

4.4 The decision

#1 TARGET — UK ACCOUNTANCY PRACTICES, 20–150 STAFF

Independent and mid-tier practices outside the Top 50. Not the Big Four, not one-person bookkeepers.

Why it wins on evidence, not vibes:

- **The pain is acute and documented.** 73% of UK firms are turning away clients for lack of staff and 74% think workloads could push people out of the profession (Advancetrack 2026). That is not a nice-to-have budget, it is a capacity crisis.
- **The timing is externally forced.** MTD for Income Tax went live on 6 April 2026 and turned annual compliance into daily work. The chasing volume went up this year and is not going back down.
- **The buyer is already convinced of the category.** 95% of mid-tier firms expect more AI, 86% have it in strategy — but only 17% believe they can assess its impact (ICAEW 2026). They want to buy and they do not know how to evaluate. A governed, auditable, measurable deployment is the answer to *their* stated problem.
- **The decision maker signs alone.** A managing partner in a 40-person practice can commit £5k without a procurement process.
- **ROI is trivially measurable.** Practice management software already counts overdue records requests, days to response and jobs by status. You do not have to invent a metric — you read theirs, before and after.
- **It matches what NOVA can do today.** Keyword retrieval over the firm's own documents with citations; email and chat channels; recurring chasing; a human approval gate on anything that leaves the building; an audit log a regulator would accept. No part of the first offer needs semantic search, a hard cost ceiling, or a field-validated channel you do not have.
- **The market is big enough and nowhere near saturated.** ~29,600–35,900 UK accounting businesses; adoption is shallow across the whole economy (~1.6 technologies per adopter).

The honest counter-argument: ICAEW is publicly flagging over-reliance on AI as "the next big threat" to firms' reputation. That is a real objection and you will hear it. It is also the best possible setup for NOVA's actual differentiator — approvals, audit, per-agent permissions and citations are the answer to exactly that fear. Lead with it rather than waiting to be hit with it.

Backup	Vertical	Why it is second, and when to switch to it
#2	Property & block management	Enormous repetitive tenant/leaseholder communication, statutory deadlines, and document volume. Scores nearly as well on workflow fit. It is second only because ROI is measured in response time rather than in a number the firm already tracks, and because the buyer is more fragmented. Switch if your first five accountancy conversations stall on professional-body caution.

Backup	Vertical	Why it is second, and when to switch to it
#3	Recruitment agencies	Very high email and CV volume, owner-operated, fast decisions, comfortable with tooling. Third because the space is crowded with point AI tools already embedded in the dominant ATS, so you compete on “why not just use the Bullhorn feature” rather than on unmet need. Switch if you find a warm route in — this vertical rewards a referral far more than cold outreach.

Two verticals to deliberately avoid first

FCA-regulated financial services scores 4.4 and should be avoided until you have three case studies and a completed security questionnaire pack. The sales cycle includes a model-risk review you cannot currently pass — no field-validated provider, no production history. **Legal firms** score 5.2 for a related reason: the work is high-value but an error is a negligence claim, and BM25 keyword retrieval without semantic search is a weak fit for case research. Both become good markets later. Neither is a first customer.

05

The first NOVA product offer

One workflow, not a platform

Do not sell “AI agents”. A managing partner does not have a budget line for AI agents and cannot tell you whether one worked. Sell the removal of one specific, named, painful, measurable job.

THE OFFER

NOVA Client Response & Records Desk

The promise: your practice stops losing partner and manager hours to chasing clients for records and answering the same routine questions — without anything reaching a client that a human did not approve, and with a log of every action your PI insurer can read.

Deployed as three named workers inside your own AWS account:

- **The Records Chaser** — knows which clients owe what and by when, drafts the chase, escalates the ones that are drifting.
- **The Client Desk** — drafts replies to routine client email from your own documented policies, and cites which document it used.
- **The Practice Librarian** — answers your staff’s “where do I find / what’s our policy on” questions from the firm’s own handbook and procedures, with citations.

5.1 Why these three, and not the exciting ones

Each maps onto a capability that is at least **TESTED** today and needs no capability NOVA does not have. None of them touches bookkeeping entries, tax computations, filings or anything a professional body would call regulated advice — which is also what makes the first sale possible.

Worker	NOVA capability it rides on	Status
Records Chaser	Governed automations + channels + approvals + audit	■ TESTED
Client Desk	Knowledge (FTS5 + citations) + channels + approvals	■ LIVE LOCAL PROVEN
Practice Librarian	Knowledge only — no outbound channel at all	■ LIVE LOCAL PROVEN

Start with the Librarian on day one

It is the only one of the three with **no outbound channel**, so it cannot embarrass anybody. It produces visible value in week one, it builds the knowledge corpus the other two depend on, and it earns the trust you need before you ask a partner to let software draft a client email. Sequence matters more than scope.

5.2 Green / Yellow / Red — decided before deployment, not after

This table goes in the contract and on the wall. It is the single most effective objection-handler you have, because it answers “what if it does something stupid” with a document rather than a reassurance. It is enforced by the compiled policy and the approvals queue, not by a prompt asking the model to be careful.

GREEN — the workforce acts alone

- Answer a staff question from the firm's own documents, with citations
- Draft a chase message into the queue for a human to release
- Summarise which clients are overdue and by how long
- Classify an inbound document by type and attach it to the right client
- Flag an approaching statutory deadline internally
- Log every action to the audit trail

Enforced by: the agent's compiled policy allow-list. Logged, not queued.

YELLOW — a human approves before it happens

- Send any message to a client (email, chat, SMS)
- Second and third chase escalations
- Anything that names a fee, a deadline commitment or a service scope
- Adding a document to the knowledge base that will shape future answers
- Any reply on a channel marked externally-visible
- Creating or changing a recurring automation


Enforced by: approval.required_for on the agent, plus per-channel approval policy. The action stops and queues; a named human releases it; both the intent and the release are audited.

RED — the workforce may not do this at all

- Tax advice, computations or anything a client could act on as advice
- Filing anything with HMRC or Companies House
- Access to the accounting ledger, bank feeds or payment systems
- Anything touching money movement, ever
- Deciding to disengage or take on a client
- Handling a complaint, a dispute or anything with a legal dimension
- Anything where the firm's professional indemnity position is in play

Enforced by: tools.deny, which compiles into the runtime's own unconditional deny list — evaluated before any bypass mode. Plus: never grant the channel, never grant the corpus, never grant the toolset.

5.3 The exact first deployment



Three to four weeks, in this order. Each stage has to be visibly working before the next is switched on.

What must be true before day one that is not true today

The Records Chaser depends on schedules actually firing, and §3.4 grades that **NOT IMPLEMENTED** — the cron ticker lives inside the runtime's gateway and no gateway process has ever been run in a NOVA deployment. Either prove that first (action 3 in §24), or scope the pilot to the Librarian and the Client Desk and trigger chases manually. **Do not sell the Chaser until you have watched a scheduled job fire.**

06

Customer ROI model

How to measure it, without inventing it

The entire commercial case rests on measuring something the customer already counts. Do not build a bespoke metric; read theirs. Every practice-management system in this market already tracks jobs by status, records requests, and days outstanding.

6.1 The measurement method

Two weeks of baseline before anything is switched on. This is non-negotiable and it is the step that gets skipped. Without it, everything afterwards is an anecdote.

Metric	Where it comes from	How NOVA contributes the “after”
Hours on the target workflow	Time recording, or a two-week tally sheet	Time recording after go-live, same categories
Records requests outstanding > 14 days	Practice management software	Same report, same filter
Median days from request to receipt	Practice management software	Same report
Median first-response time to client email	Mail system reporting	Same report
Tasks completed by the workforce	n/a	NOVA Work screen — durable task board
Automation rate (acted alone ÷ total)	n/a	NOVA audit log: committed actions with no approval record
Human escalation rate	n/a	NOVA Approvals: queued ÷ total
Error / correction rate	Manual review of a sample	Sample 30 outputs weekly and have a manager grade them. Do this honestly — it is the number the customer will ask about in month three.
Cost per workflow instance	(loaded hourly rate × hours) ÷ instances	Same calculation with the after-hours
Revenue protected	Fees at risk from capacity turn-away	Only claim this if the firm can name the clients it turned away

6.2 The calculation

Monthly labour value recovered

= (baseline hours – after hours) × loaded hourly rate

Loaded hourly rate

= (salary × 1.30 for employer NI, pension, overhead) ÷ 1,650 working hours

Net monthly benefit

= labour value recovered – NOVA monthly fee – AWS + model spend

Payback period (months)


= implementation fee ÷ net monthly benefit

Use the customer's own salary figures. If they will not share them, UK market data for an accounts assistant sits around £24k–£27k depending on source (ONS ASHE-derived and commercial salary trackers, 2026) — but their number beats any published one, and asking for it is itself a qualifying question.

6.3 An illustrative scenario

Everything in this scenario is an assumption

The hours below are **invented for the purpose of showing the arithmetic**. NOVA has no customers and therefore no results. Use this structure with the customer's own baseline numbers filled in; never present these figures as evidence, and never put them on a slide without this warning attached.



ILLUSTRATIVE ASSUMPTION. A hypothetical 30-person practice. Not measured, not a benchmark, not a promise.

Line	Illustrative value	Where a real number would come from
Baseline hours/month on the five workflows	138	Two-week tally × 2
Assumed hours after	42	Same tally, post go-live
Hours recovered	96	Subtraction
Assumed loaded rate	£21/hour	Firm's own payroll ÷ 1,650
Labour value recovered	£2,016/month	96 × £21
Less NOVA monthly fee	–£950	\$7
Less AWS + model spend	–£250 (assumed)	Actual bill — see §18
Net monthly benefit	£816	
Implementation fee	£4,500	\$7
Payback	~5.5 months	4,500 ÷ 816

ILLUSTRATIVE ASSUMPTION throughout. The purpose of this table is the method and the shape of the argument, not the numbers.

The honest thing to say in the room

“I can't tell you what you'll save, because you'd be our first customer in this sector and I'm not going to quote you someone else's number. What I will do is measure your baseline for two weeks before we switch anything on, and at day 60 we look at your own report together. If it hasn't moved, you've paid for a pilot and you stop.”
— This wins more deals than a fabricated case study, and it is the only version you can actually defend.

07 Pricing

What to charge, and what never to charge for

7.1 What the UK market is paying

Published 2026 benchmarks for UK and international AI automation work cluster as below. Treat these as the shape of the market, not as a price list — most are published by agencies with an interest in the number.

Segment	Implementation	Monthly	Source type
Single scoped workflow, UK	£1,000–£5,000	£200–£800	UK agency pricing surveys, 2026
Connected multi-workflow build, UK	£4,000–£12,000	—	UK agency pricing surveys, 2026
Small business, 2–3 workflows	\$5,000–\$25,000	\$1,000–\$3,500/mo	International agency benchmarks, 2026
Mid-market AI stack + reporting	—	\$4,000–\$10,000/mo	International agency benchmarks, 2026
Custom agent development	from \$10,000, \$50,000+ complex	—	International agency benchmarks, 2026
UK marketing agency AI tooling	—	£350–£1,100 per seat/mo	UK agency reporting, 2026

Sources listed in §25. These are self-published agency benchmarks and should be read as directional. The prevailing 2026 structure is consistent across all of them: setup fee + monthly retainer + usage.

7.2 Recommended pricing

FOUNDING CUSTOMER

£4,500 implementation · **£950 / month**

- Customers 1–3 only. Price is explicitly time-limited and named as such.
- 12-month term, price locked for 24 months.
- One workflow, up to three agents, one knowledge corpus, two channels.
- In exchange — and this is written into the agreement — a named case study, a reference call, and access to their baseline numbers.
- Customer pays AWS and model costs directly, in their own account.
- Fortnightly review call for the first 90 days.

STANDARD SMB

£7,500 implementation · **£1,450 / month**

- 20–100 staff. The main line once you have two case studies.
- One workflow area, up to five agents, two corpora, three channels.
- 12-month term, monthly thereafter, 60 days' notice.
- Quarterly optimisation review included.
- Additional workflow: £3,500 setup, +£450/month.
- Bespoke integration: quoted separately, never bundled.

MID-MARKET**£15,000–£25,000** implementation · **£2,800–£4,500 / month**

- 100–400 staff, or multi-site.
- Unlimited agents within an agreed policy envelope; up to five corpora.
- Named delivery lead, monthly governance review, quarterly security review.
- Priority response SLA (business hours).
- Assumes a security questionnaire and a DPIA — budget two to six weeks for it.

ENTERPRISE**£40,000+** implementation · **£6,000+ / month**

- Do not chase this yet. You have no production history, no SOC 2, no on-call rota and no field validation. You will lose the deal in the security review and burn six months doing it.
- Revisit after three referenceable customers and a completed security pack.
- When you do: annual prepay, custom DPA, named architect, defined RTO/RPO.

7.3 Why the customer pays

Tie every pound to a business outcome and never to consumption. The sentence that sells this tier is: **“the monthly fee is less than a third of one part-time administrator, and it is the only administrator whose every action is logged.”**

The fee buys	Not	Because
A workflow removed from partners' desks	“Access to AI”	They can get access to AI for £20/seat. They cannot get this.
Governance they can show an insurer or a regulator	A number of prompts	Approvals, per-agent permissions, and a write-ahead audit log are the differentiator and the hardest thing to rebuild.
Deployment into their own account	A seat on your platform	Their data never leaves their AWS account. For this buyer that is worth more than any feature.
Someone accountable for it working	Software	At this size, the customer is buying an outcome with a person behind it.

7.4 What not to do

PRICING MISTAKES THAT WILL COST YOU THE BUSINESS

- **Never price per prompt, per token or per message.** It makes the customer ration usage, which kills adoption, which kills the renewal. It also makes your revenue a function of model pricing you do not control.
- **Never include unlimited model usage in a fixed fee.** You have no hard spending ceiling — \$3.8 — and one runaway loop on an expensive model can exceed a month's fee. The customer pays their own AWS and model bill, in their own account. This is not a concession; it is the architecture, and it is also a selling point.
- **Never promise unlimited custom development.** “And we'll build whatever integrations you need” has destroyed more early-stage software companies than any competitor. Each integration is quoted.
- **Never discount without taking something back.** A discount in exchange for a case study, a reference, a longer term or a faster decision is strategy. A discount because they pushed is a signal that the price was invented, and they will push again at renewal.

- **Never let the implementation fee disappear into the monthly.** It funds the real cost of delivery and it qualifies the buyer. A customer who will not pay to start will not stay.
- **Never quote before discovery.** You cannot price a workflow you have not watched someone do.

08

Getting the first UK customer

Where they are, and how to reach them

This section is deliberately concrete. Named sources, named criteria, named sequence. A first customer comes from about 100 well-chosen conversations, not from a marketing funnel you do not yet have traffic for.

8.1 Where the prospects actually are

Listed by how well each source lets you **filter**, which matters far more than how many records it holds. I have marked what I have verified exists versus what you should check yourself before relying on it.

Source	What it gives you	Verified?
ICAEW / ACCA / AAT member firm directories	Firm name, location, services, professional body. The highest-quality starting point because membership itself is a qualifier.	Directories exist and are public. Check current search filters yourself.
Companies House	Free API and bulk data. Filter by SIC code 69201/69202 (accounting, bookkeeping and auditing), incorporation date, filing history, registered office. Gives you officer names — i.e. the partners.	Public and free. SIC codes are the standard classification.
LinkedIn Sales Navigator	The main working tool. Filter by industry (Accounting), headcount 11–200, geography, and seniority (Partner / Managing Partner / Practice Manager). Shows recent hiring, which signals capacity pressure.	Paid. Budget ~£70–£100/month — the single best-value spend in this plan.
Accountancy trade press (AccountingWEB, Accountancy Age, Accountancy Today)	Named practitioners talking publicly about capacity, MTD and AI. A partner who has written about being overloaded is a warm lead.	Publications verified active in 2026.
Accountex and regional practice events	The main UK accountancy technology event. Also regional ICAEW branch meetings, which are smaller and far easier to actually meet people at.	Accountex London ran in 2026. Check dates and cost yourself.
Practice software ecosystems (Xero, Sage, IRIS, TaxCalc partner lists)	Firms already comfortable buying software. Partner directories are public.	Directories exist. Terms of use vary — check before scraping.
Referrals from your own network	Every business has an accountant. Ask yours. Ask your solicitor's. This is the highest-conversion channel by a wide margin and costs nothing.	n/a
Local business networks (chambers of commerce, BNI, Federation of Small Businesses)	Accountants are heavily over-represented in these, because that is where they find clients. So it is where you find them.	n/a
Google Maps / local directories	Coverage of independent practices by town. Useful for geographic completeness, weak on firm size.	n/a
Clutch, Upwork, agency directories	Largely irrelevant for this vertical. Useful if you switch to #3 (recruitment) or digital agencies.	n/a

8.2 Prospect criteria

Must have	Strong signal	Disqualify
20–150 staff	Recently hired or is advertising for admin/support roles	Under 10 staff — no budget, owner does everything
UK-based, single or few offices	Public comment about capacity or MTD workload	Top 50 firm — procurement, InfoSec, and an internal team
Cloud practice software already (Xero / Sage / IRIS / Karbon)	Has a practice manager or operations manager (a named process owner)	No cloud software at all — the integration story collapses

Must have	Strong signal	Disqualify
An identifiable partner or MD you can reach	Website mentions client service or responsiveness as a differentiator	Already has an internal AI/automation lead — you become a bake-off
Client base of 100+ (so chasing volume is real)	Growing: new office, new hires, acquisition	Currently in a merger — nothing gets decided

The qualifying question that saves you the most time

“Roughly how many hours a week does your team spend chasing clients for records?” If they answer with a number, they have thought about it, it hurts, and you have a live opportunity. If they say “no idea, not much”, politely move on — you cannot sell measurable ROI to someone who does not measure. Ask it in the first five minutes of every call.

09

The first 100 prospects

A list, a score, and a 30-day plan

9.1 The tracker

One spreadsheet. Build it by hand — the research *is* the qualification, and outsourcing it produces a list you cannot sell to.

Column	What goes in it
Company	Legal name from Companies House
Website	
Industry / niche	e.g. 'owner-managed businesses', 'contractors', 'dental'
Employees	LinkedIn, sanity-checked against the website's team page
Decision maker	Named human
Role	Managing Partner / Partner / Practice Manager
LinkedIn URL	
Email	Firm's published address until you have a reason to use a personal one
Pain hypothesis	One sentence, specific to them — this is the whole email
Existing software	Xero / Sage / IRIS / Karbon / unknown
AI maturity	None / experimenting / tooling in place
Priority	A / B / C from the score below
Contacted	Date + which touch
Response	None / negative / neutral / positive
Meeting	Date
Proposal	Date + value
Won / Lost	+ one-line reason. This column is how you learn.


9.2 Lead scoring

Points	Signal	
+3	20–150 staff	
+3	Named, reachable partner or MD	
+2	Cloud practice software confirmed	
+2	Hiring admin/support, or publicly discussing capacity	
+2	Warm route in (referral, mutual connection, met at an event)	
+1	Active on LinkedIn — they will see your content	
+1	Website emphasises responsiveness or client service	
−2	No cloud software	
−3	Top 50 firm, or an internal AI lead	
−5	Under 10 staff	
Band	Score	Treatment
A	9+	Personal research, multi-touch, phone. ~25 of the 100.

Band	Score	Treatment
B	5–8	Templated-but-specific email plus LinkedIn. ~50.
C	<5	One email. Do not spend more. ~25.

9.3 The 30-day plan

	Focus	Daily target	End-of-week gate
Week 1	Research and list building. No outreach at all — resist this.	20 fully-researched prospects/day incl. a specific pain hypothesis	100 rows complete and scored. 25 A-band identified.
Week 2	First-touch outreach and follow-up. Start with B-band to rehearse the message before you spend your A-band.	25 first-touch emails + 10 LinkedIn connections + 5 calls	All 100 contacted once. At least 40 second touches sent.
Week 3	Conversations. Every reply gets a call offer within two hours.	5 calls attempted, 2 discovery calls held, all follow-ups within 24h	8–12 discovery calls held. 4–6 demos booked.
Week 4	Demos and pilot proposals. Proposal within 24 hours of every demo, while it is still warm.	2 demos/day, proposal out same day	2–3 pilot proposals live. One verbal commitment.



ILLUSTRATIVE PLANNING TARGET. Set from general cold-outreach norms because NOVA has no conversion history of its own. Use it to size activity, and replace every number with your own after the first 100.

If 100 prospects produce nothing

That is information, not failure, and it means one of exactly three things: the message is wrong, the vertical is wrong, or the offer is wrong. Diagnose in that order — message is cheapest to change and is wrong most often. Do **not** respond by sending another 100 of the same email; that tells you nothing you did not already know.

10

Marketing and advertising

What to do, and what to skip

10.1 Organic — where nearly all your effort goes

Channel	What to publish	Cadence	Why it works here
LinkedIn, founder account	Not the company page — nobody follows company pages. Build a personal audience of UK accountants.	3×/week	Your buyer is genuinely active here, and a partner will read a founder before a brand.
The governance angle	“What happens when an AI agent gets it wrong, and who finds out”. Screenshots of a real approvals queue and a real audit log.	1×/week	ICAEW is publicly flagging over-reliance on AI as a reputational threat. You are selling the answer to a fear the profession has already named.
Before/after workflow posts	One workflow, the hours it takes, what it looks like governed. No product screenshots for the first three lines.	1×/week	Concrete beats conceptual with this audience every time.
Short video (60–90s)	Screen recording of the Control Center doing one real thing end to end. No music, no motion graphics.	2×/month	Proves it exists. Most of your competition at this stage is a landing page.
Technical architecture posts	How single-tenant BYOC deployment works. IAM boundaries. Why no secret is stored.	1×/month	Reaches the technically-minded partner and anyone doing due diligence on you.
Customer ROI stories	Only once you have one. Named firm, their numbers, their quote.	when earned	This is the single highest-value marketing asset you will ever produce — which is why the founding-customer discount buys it contractually.
Open-source credibility	Be straight that NOVA governs the open-source Hermes runtime. Contribute upstream where you can.	ongoing	Honesty here is a differentiator. It also pre-empts the discovery, which is far worse if the buyer makes it themselves.

10.2 Direct sales — where the first customer actually comes from

Motion	Volume	Notes
Cold email, researched	25/day in week 2, then 10/day sustained	Four to six sentences. Their specific pain, one sentence on the mechanism, one ask. No attachment, no deck, no calendar link in the first email.
LinkedIn outreach	10 connections/day	Connect, then nothing for a week. Engage with their posts. Message only when you have a reason. Pitching in the connection request destroys the account.
Phone	5 attempts/day on A-band only	Still works in this market and almost nobody does it. Call the practice, ask for the practice manager. 20 seconds on why you called, then ask for 15 minutes.
Partner referrals	ongoing	Bookkeepers, fractional FDs, practice-software consultants, IT MSPs serving accountancy. They already have the relationship and no competing product. \$15.
Events	1–2 per quarter	Regional ICAEW branch events over big trade shows. Ten real conversations beats 400 badge scans, and costs a hundredth as much.

10.3 Paid advertising — the recommendation is no

DO NOT RUN PAID ADS UNTIL CUSTOMER THREE

This is a considered recommendation, not caution for its own sake. Paid acquisition converts attention into customers. You currently have no proven conversion, no case study, no measured ROI claim, and no landing page that has ever converted anyone. Spending on ads now buys traffic to a page that cannot close, and you will read the resulting zero as “ads do not work” when what it actually measured was the offer.

There is also an arithmetic problem specific to this market. LinkedIn Ads to a senior UK professional-services audience is among the most expensive inventory there is. At a realistic cost per qualified lead in that channel, and a first-deal value of roughly £4.5k setup plus £11.4k of first-year recurring, you would need a conversion rate from cold ad click to signed pilot that nobody achieves without an established brand and a case study. The maths does not close, and no budget fixes it.

When each channel becomes worth testing:

- **Google Search** — first, and only for high-intent phrases people actually type when they have already decided (“ai automation for accountancy practice uk”). Volume will be tiny, which is the point: tiny and intent-loaded is exactly what you want. Start at **£300–£500/month** for one month and judge on qualified conversations, not clicks.
- **LinkedIn Ads** — after two case studies, and then only as **retargeting** of people who already visited. Cold LinkedIn prospecting ads are for companies with brand recognition. Budget **£500–£1,000/month** when you get there.
- **Retargeting generally** — the first paid spend that makes sense, because it is the only one aimed at people who have already shown intent. Even so it needs meaningful traffic first.
- **Meta Ads** — **no**. Wrong audience, wrong intent, wrong buying mode for a £15k+ annual B2B professional-services purchase.

Where that budget should go instead

LinkedIn Sales Navigator (~£70–£100/month) plus an email-verification tool plus your own time. For the price of one month of LinkedIn Ads you can research and contact several hundred named firms with a message written for each. At this stage that converts and advertising does not.

11

The NOVA message

What to say, in every length

Written against what NOVA actually does. Nothing here claims field validation, semantic search, spending caps, or a channel that has not been connected.

ONE-LINE PITCH

NOVA puts a governed AI workforce inside your business — in your own cloud, connected to your systems, and answerable to your rules.

30-SECOND PITCH

“Most firms have tried AI and got a chat assistant — it drafts things, and the admin comes straight back. NOVA is different: it deploys named AI workers *inside your own AWS account* that read your documents, work your channels, and handle recurring jobs like chasing clients for records. Every worker has a job description you write, permissions you set, and anything consequential stops for a human to approve. Every action is logged with who authorised it. Your data never leaves your account.”

Roughly 75 words. Time it — 30 seconds is shorter than people think.

60-SECOND PITCH

“A 40-person practice loses something like two working days a week to chasing clients for records and answering the same routine questions. You’ve probably tried AI for it. It helped with drafting and then the admin came straight back, because the assistant can’t see your systems and can’t be trusted to act.

NOVA deploys AI workers into your own AWS account — your infrastructure, your data, your encryption key. Each one has a job description you write and permissions you set. The Records Chaser knows who owes what and drafts the chase. The Client Desk answers routine client email from your own documented policies and cites which document it used. The Practice Librarian answers your team’s questions from your handbook.

Anything that reaches a client stops and waits for a human to release it. Every action — including every refusal — is written to an audit log with the person who authorised it. That’s the part your PI insurer will care about.

We start with a 60-day paid pilot. Two weeks measuring your baseline first, so at the end we’re looking at your numbers, not mine.”

Asset	Copy
Website headline	Your AI workforce. Your cloud. Your rules.
Website subheadline	NOVA deploys governed AI workers into your own AWS account — connected to your documents, your channels and your recurring work, with human approval on anything that matters and a full audit trail on everything else.
LinkedIn company description	NOVA is an enterprise AI workforce platform for professional services firms. We deploy governed AI agents into your own cloud account, connected to your knowledge, systems and communication channels — with per-agent permissions, human approval on consequential actions, and a complete audit trail. Built on the open-source Hermes agent runtime. Your data never leaves your infrastructure.

Asset	Copy
Founder explanation (for "so what do you do?")	"I build the layer that makes AI agents safe to actually deploy in a business. The agent technology is open source and it's good. What's missing is everything a real company needs before it lets software act on its behalf — permissions, approvals, an audit trail, and running it in your own infrastructure instead of someone else's. That's what I build. Right now I'm working with accountancy practices, because they're drowning and they're regulated, which is the exact combination this is for."
Sales presentation opening	"Before I show you anything — roughly how many hours a week does your team lose to chasing clients for records? [listen] And what happens today when a client just doesn't respond? [listen] Right. Let me show you what that looks like when it's handled, and I'll be specific about what this can't do as well as what it can."

11.1 Words that are banned, and what to say instead


Never say	Say instead	Because
Revolutionary / game-changing / cutting-edge / next-generation	Nothing. Delete the sentence.	These signal that you have no specific claim. A partner reads them as noise.
"Fully autonomous"	"Acts alone on what you've approved it to act alone on"	Autonomy is the fear, not the benefit, for this buyer.
"Replaces your admin team"	"Gives your team back the hours they spend chasing"	The buyer has a staffing shortage, not a staffing surplus. 95% of AI-using SMEs report no workforce reduction (BCC 2026) — so it is also untrue.
"Understands your documents"	"Searches your documents and cites what it used"	It is BM25 keyword retrieval. Say what it is; the citation is the actual benefit.
"Enterprise-grade security"	Name the controls: own account, own KMS key, no inbound ports, SSM-only access, IAM permissions boundary, audit log.	Specifics are checkable. The phrase is not.
"It just works"	"Here is the 60-day pilot and here is what we measure"	Nobody at this price point believes the first one.

12

Website structure

Ten pages, and what goes on each

12.1 The homepage



The order matters. Security sits above integrations because for this buyer it is the first objection, not the last.

Section	What it must contain	What it must not
Hero	The headline from §11, a one-line subhead, and exactly one CTA: "Book a 15-minute call". A still of the real Control Center.	A carousel. A second CTA. A stock photo of a robot.
Problem	The specific hours: chasing records, routine client email, internal questions. Written so a partner recognises their own week.	Statistics about the AI market. They do not care.
AI Workforce	The three named workers, one card each, with the one sentence each does.	A feature grid of twelve capabilities.
How NOVA works	The four-layer diagram from §2, simplified. Name Hermes honestly in one line.	Architecture depth. Link to a separate page for that.
Security	Your AWS account. Your KMS key. No inbound ports. SSM-only access. No credential stored by us. Audit log of every action.	"Bank-level security". A padlock icon. Compliance badges you do not hold.
Integrations	Honest categories: 22 channel platforms available, 65 MCP servers, your own documents. A note that each is enabled per agent.	A logo wall . Zero channels are field-validated — a Slack logo is a claim you cannot currently support.
ROI	The measurement method from §6. "We measure your baseline for two weeks before we switch anything on."	Invented percentages. Any figure you cannot attribute.
Use cases	One page per worker, each with the Green/Yellow/Red table.	Generic 'AI for business' content.
Demo	A 90-second screen recording of one real thing, end to end.	A motion-graphics product film. It reads as vapourware.
CTA	Calendar link. Name, firm, size, one line on the workflow.	A 9-field form. A chatbot.

12.2 The other nine pages

Page	Job it does
How it works	For the person who has to explain it internally. The full architecture, the deployment model, what apply does, what an agent is.
AI Workforce	The capability catalogue in customer language. Be graded — an "available / in pilot" column earns more trust than it costs.
Governance	Your strongest page. Build it first after the homepage. Permissions, approvals, audit, Green/Yellow/Red, screenshots of a real audit log.

Page	Job it does
Integrations	Categories and mechanism, not logos. Explain that a connection is granted to a named agent and revoked in one click.
Security	For the buyer's IT person or MSP. Deployment topology, IAM model, encryption, where data lives, what you can and cannot see. Offer to answer a security questionnaire — most competitors at this size will not.
Industries	One page: accountancy. Ten thin pages are worse than one page that is obviously written by someone who understands the work.
Pricing	Publish the implementation fee and the monthly. Your buyer is a professional-services firm; they respect transparent pricing and “contact us” loses you the ones who would have bought.
Case studies	Empty until you have one. An honest “our first pilots are running now” beats three invented logos, which is the fastest way to lose a partner's trust permanently.
Contact / demo	Calendar. Four fields. Nothing else.

13

The 15-minute sales demo

A storyline, not a tour

The demo's job is to make a partner believe two things: that this removes a job they hate, and that it cannot embarrass them. Everything that does not serve one of those is cut, including most of the architecture you are proud of.

Three rules for the whole 15 minutes

1. Use their language. “Records request”, not “task”. “Client”, not “end user”. “Partner sign-off”, not “approval workflow”. **2. Never say the word “agent” in the first eight minutes.** Say “a worker called the Records Chaser”. **3. Name one limitation out loud, unprompted.** It buys more credibility than any feature, and they will find it anyway.

MINUTE 0–2 | The business problem

- Screen off. You are talking, not showing.
- “Before I show you anything — how many hours a week does your team lose to chasing records?” Then listen properly.
- Play their number back with the annual cost attached. “So that's roughly £X a year of partner and manager time on chasing.”
- “And the bit I hear most is that it isn't the time, it's that it never ends. Is that fair?”

MINUTE 2–5 | The Control Center

- Open on **Overview**. Their firm's name and their brand colour, because white-labelling is real and it lands instantly.
- “This is your workforce. Three workers. This is what each one has done this week.”
- Click into **Work**. Real tasks, real states.
- Do not explain the architecture. They have not asked.

MINUTE 5–8 | The workforce itself

- Open the **Records Chaser** profile.
- **Soul** tab: “this is its job description — in plain English, and you can edit it.” Change a line and save it in front of them. This is the moment most demos are won.
- **Permissions** tab: “this is everything it's allowed to touch. Nothing else is on this list, so nothing else can happen.”
- “And if you want a fourth worker, you create one here.” Do not actually create it — you will burn 90 seconds.

MINUTE 8–10 | Knowledge and connections

- **Knowledge**: upload one of *their* documents live if they have sent one. Otherwise a realistic sample handbook.
- Ask a question against it. Show the answer **with the citation**.

- “It tells you which document it used. If you change the document, the answer changes. It isn't making it up from the internet.”
- **Say the limitation here:** “This is keyword search over your documents, not the fuzzy semantic kind. It's excellent when your documents use your firm's own vocabulary, which yours do. I'd rather tell you that now than have you find it in week three.”
- **Channels:** show the list and the per-agent grant. Do **not** claim any specific platform is live unless you have connected it for them.

MINUTE 10–12 | Governance — the part that closes it

- **Approvals.** “This client email is drafted and it is *not* going anywhere until someone here releases it.” Release one.
- **Activity.** “Every refusal and every escalation, with who authorised it and when.”
- Show the Green/Yellow/Red table. “We agree this before we deploy anything, and it's in the contract.”
- “Your PI insurer and your professional body are going to ask what happens when it gets something wrong. This is the answer.”

MINUTE 12–14 | Recurring work

- **Automations.** “The Monday-morning chase runs on a schedule. Here it is, here's when it next runs, here's what it did last time.”
- “And it can't be given a new standing instruction without going through the same permission checks as everything else.”
- **Say the limitation here too** if the scheduler is not yet field-proven in your deployment: “Scheduling is the newest part of this and I'll be proving it in your environment in week one of the pilot rather than claiming it today.”

MINUTE 14–15 | ROI and the ask

- Back to their number from minute one.
- “You said about 30 hours a week. We measure that properly for two weeks before we switch anything on, so at day 60 we're looking at your report, not my slide.”
- **The ask:** “I'd like to run a 60-day pilot. £4,500 to set it up, £950 a month. You pay AWS directly — it'll be a few hundred a month in your own account. If the numbers haven't moved at day 60, we stop and you've paid for a pilot.”
- Then stop talking.

What to cut when you run out of time — and you will

Cut minutes 12–14 (automations) before anything else. Never cut 10–12 (governance): it is the section that differentiates you from every other person who has emailed this partner about AI this year.

14

The first customer pilot

60 days, with a defined end

A pilot with no end date becomes a free deployment. This one has a fixed length, a fixed scope, an agreed definition of success written before it starts, and a decision at the end.

WEEK 0	Discovery and baseline	Watch the workflow being done. Two-week baseline measurement starts now, before anything is installed. Agree Green/Yellow/Red. Name the success criteria and the one person on their side who owns it.
WEEK 1	Infrastructure and integration	AWS account access, Terraform apply, image push, tenant bundle, principals, model credentials. Knowledge corpus loaded from their handbook. Nothing client-facing yet.
WEEK 2	Agent configuration	The Practice Librarian goes live to internal staff only. Personas written with them, not for them. Policy compiled and verified by attempting a denied action in front of them.
WEEK 3	Controlled deployment	Client Desk in draft-only mode — every message queued for approval. Records Chaser triggered manually, not scheduled, until the scheduler is proven in their environment.
WEEK 4–6	Measurement and widening	Approval rate should be falling as the drafts get better. Widen Green as trust earns it. Weekly 30-minute review with their owner. Sample 30 outputs a week and grade them honestly.
DAY 60	ROI review and decision	Their numbers, their report, side by side with the baseline. Three outcomes: convert to standard, extend once with a named reason, or stop. Agree which before you walk in.

14.1 Success criteria

These are proposed pilot targets, not predictions

NOVA has no customers, so no target here is derived from experience. Set them *with* the customer in week 0 against their own baseline, write them down, and sign them. A target you set unilaterally is one they will dispute at day 60.

Category	Proposed measure	Proposed target	Why this one
Hours saved	Hours/month on the target workflow vs baseline	–40% or better	The headline. If this moves, nothing else matters much.
Response time	Median first response to routine client email	–50%	Visible to their clients, so it is felt as well as measured.
Workflow completion	Records requests closed within 14 days	+25pp	Their existing report, unchanged — no argument about methodology.
Escalation rate	Share of actions needing human release	Falling week on week	The trajectory matters more than the level. Flat means it is not learning the firm's voice.
Error rate	Outputs a manager marks wrong, from a 30/week sample	<5% by week 4	Measure it honestly or it will be measured for you by a client complaint.
Adoption	Staff using the Librarian weekly	>60% of fee earners	Unused software does not renew, however good the hours number is.
Cost	Total AWS + model spend	Within agreed ceiling	There is no hard cap in the product — see §3.8. Watch it manually and weekly.

14.2 What kills a pilot

Failure	Prevention
No baseline measured	Week 0 is not optional. Refuse to start without it.
Scope grew	One workflow. Everything else goes on a written 'phase 2' list — which also becomes the upsell.
No owner on their side	Name one person in week 0. If they cannot name one, they are not ready.
First client-facing output was wrong and nobody saw it	Draft-only for the whole of week 3. No exceptions, however good it looks.
Champion left / got busy	Two contacts from day one, and the partner in the weekly review even if they say they do not need to be.
The scheduler never fired	Prove it in week 1 in <i>their</i> environment, before anything depends on it.

15

The UK sales partner model

Commission-only, structured properly

Published 2026 B2B benchmarks put new-business SaaS commission around 8–12% of new ARR, renewals at 2–8%, and pure-commission independent representatives materially higher — commonly quoted in the 15–30% range to compensate for carrying all the risk. The structure below sits inside those norms while protecting cash flow and, more importantly, protecting the customer relationship.

15.1 Recommended structure

Term	Recommendation	Reasoning
Status	Self-employed contractor, written agreement. Not an employee.	No PAYE, no notice period, no employment risk. Have a UK employment solicitor check it — misclassification is expensive.
Base	None initially. Revisit at customer five.	You cannot fund a salary pre-revenue, and a commission-only partner who will not take the risk does not believe in the product.
Implementation commission	20% of the setup fee	≈£900 on a £4,500 founding deal. Paid on cleared funds , never on signature.
Year-1 recurring	15% of monthly revenue, months 1–12	≈£143/month on a £950 deal. Sits in the independent-rep band and rewards landing a customer who stays.
Year-2 recurring	7% for as long as the customer is live <i>and</i> the partner still owns the relationship	Inside the 2–8% renewal norm. Drops to 0% if account management moves to you.
Accelerator	Add 5pp on implementation after the 4th customer in a 12-month window	Rewards a repeatable motion rather than one lucky deal.
Payment timing	Monthly in arrears, within 14 days of the customer's cleared payment	Never pay commission on revenue you have not banked.
Clawback	Full clawback if the customer cancels or refunds within 90 days	Stops a partner selling a deal that was never going to stick.
Lead attribution	Written register. A lead is the partner's for 90 days from first registration; inbound leads to you are yours unless you assign them.	Attribution disputes end partnerships. Decide it on day one, in writing.
Customer ownership	Yours, absolutely. Contract, data, account and relationship are with your company.	Non-negotiable. A partner who wants to own the customer is a reseller, which is a different agreement entirely.
Territory	UK. Optionally a named vertical.	Do not fragment a market this size geographically.
Exclusivity	None for the first 6 months. Then optional, conditional on a quota.	Exclusivity given before performance is proven is how you lose a market to one person's calendar.
Termination	30 days either side. Commission on already-closed customers continues for 6 months after termination, then stops.	Fair, finite, and it keeps the exit clean.
Confidentiality	Mutual NDA. Explicitly covers customer names, the repository, pricing floors and the capability gaps in §23.	They will see the honest limitations. That is deliberate and it must be protected.
Authority to quote	May quote the published tiers only. Any discount, any custom scope, any term change needs your written approval.	A partner discounting to close costs you the price anchor for every future deal.
No unauthorised technical promises	Written list of what may and may not be claimed, derived from §3 and §23. Breach is a termination event.	The single most important clause. A commission-only partner is structurally incentivised to promise anything. See the box below.

The capability schedule — attach it to the agreement

Give the partner a one-page annex listing exactly what may be claimed: governed agents, per-agent permissions, human approval, audit log, deployment in the customer's own AWS account, keyword document search with citations, 22 available channel platforms. And exactly what may **not**: any named channel as “live”, semantic or AI search, a hard spending cap, scheduled execution, multi-tenancy, any compliance certification, and any customer result. Update it every time §3 moves a rung. A promise your partner makes is a promise you have to deliver.

15.2 When to hire a full-time salesperson

Do NOT hire until all of these are true	Why
3+ paying customers, at least 2 referenceable	You are selling a proven thing, not testing whether it sells.
A repeatable sales process — the same demo and objections work twice	Otherwise you are hiring someone to discover your process for you, at salary.
The funder has personally closed at least 3	You cannot hire out a motion you have not run. You also cannot coach it or interview for it.
Deployment is repeatable — a second customer deployed without heroics	Sales that outrun delivery burn the reputation you are selling.
12 months of the loaded cost in the bank	UK B2B sales hires take 3–6 months to produce. Loaded cost is salary + ~30% for NI, pension and overhead.
A pipeline they can work from day one	A new hire who must also build the list will take nine months, not six.

Commission benchmarks: 2026 B2B and SaaS commission surveys (Everstage, CaptivateIQ, Fullcast) — see §25. Interpret as market norms, not as a published standard.

16

Commercial best practices

The rules, in five areas

PRODUCT

- **Sell outcomes, not agents.** “Chasing handled” has a budget line. “AI agents” does not.
- **Start narrow and stay narrow.** One workflow, three workers. Everything else goes on a written phase-2 list — which is also your upsell pipeline.
- **Prove ROI with their numbers.** Never present a figure you generated.
- **Do not build ahead of a customer.** Every hour spent on a feature nobody has asked for is an hour not spent on field validation, which is the actual blocker.
- **Never promise a capability the repository cannot support.** §3 is the source of truth. When it moves a rung, the sales material moves with it — not before.
- **Say one limitation out loud in every sales conversation.** It is the cheapest credibility available and it inoculates you against discovery later.

SECURITY

- **Least privilege, always.** The runtime IAM role carries no customer-data permissions and sits under a permissions boundary. Keep it that way even when a customer offers broader access to make an integration easier.
- **Customer-owned credentials.** NOVA never stores a secret. The profile `.env` is the customer's file and survives apply. Do not build a feature that changes this.
- **No secrets in images, ever.** Eight checks enforce it in the build. If a customer asks you to bake a key in, the answer is no.
- **One tenant, one deployment.** Until multi-tenancy exists — and it does not — this is the isolation guarantee. Do not co-locate two customers.
- **Audit everything consequential.** Intent before, outcome after, with the authenticated human. This is the product's spine, not a feature.
- **Human approval on anything that leaves the building.** Widen Green only when measured accuracy earns it, never because it is slow.
- **Never grant an agent broad access to the customer's AWS account.** Each integration is a separate role assumed with an ExternalId.

SALES

- **Do not lead with technology.** Eight minutes before the word “agent”.
- **Quantify the pain in the first five minutes** or disqualify. A prospect who cannot estimate the hours cannot evaluate the result.
- **Sell to the partner, not the office manager.** The office manager is your champion and cannot sign.
- **Always a paid pilot.** Free pilots do not get resourced, do not get measured, and do not convert.
- **Build the case study contractually,** in the founding-customer agreement. Asking afterwards gets a polite no.
- **Ask for referrals at the ROI review,** when the number is in front of them — not at renewal, when they are thinking about cost.
- **Write down every lost deal's reason.** Ten losses tell you more than your first win does.

PRICING

- **Price the outcome.** Benchmark against the loaded cost of the person doing the work today, not against a software subscription.
- **Protect the margin.** Implementation fee covers real delivery cost. If it does not, the price is wrong.
- **Keep implementation and recurring separate.** They fund different things and they qualify differently.
- **Charge for every custom integration.** Separately, and quoted after scoping.
- **Never promise unlimited usage.** There is no hard cost ceiling in the product. The customer's own AWS account is the correct answer and a good one.
- **Raise prices after the second customer.** The founding price is a time-limited instrument, and it must be named as one from the first email.

DELIVERY

- **Discovery before quoting.** Watch the work being done, in person if you can.
- **Architecture and security review before deployment.** Their IT person or MSP will be involved eventually — better on your terms in week one.
- **Deploy with Terraform, never by hand.** A hand-built deployment is one you cannot reproduce, and the second customer is where that bill arrives.
- **Validate before handover.** Health, audit, isolation, permissions, restart, restore. \$19.
- **Train the humans, not just the software.** The approvals queue is a new habit for somebody. Sit with them.
- **Monitor from day one.** CloudWatch, the audit log, and a weekly look at the model bill.
- **Optimise monthly for the first quarter.** Personas and Green/Yellow/Red lines both need tuning once real work flows through them.

17


Deploying NOVA to AWS

Step by step, for a non-engineer

This assumes you understand business but not cloud infrastructure. Every technical term is explained the first time it appears. Work through the steps in order — several of them depend on the one before.

17.1 What you are actually doing

You are putting a software package onto a computer you rent from Amazon, inside the customer's own Amazon account, with a disk that survives restarts, permissions that are deliberately tiny, no way in from the internet, and a record of everything that happens. That is the whole thing.



17.2 Every term, explained once

Term	What it actually is
Docker	A way of packaging software so it runs the same everywhere. A Docker image is the package; a container is that package running. NOVA ships as one image, 304 MB.
ECR	Elastic Container Registry . Amazon's private warehouse for Docker images. You upload the NOVA image here; the server downloads it from here.
EC2	Elastic Compute Cloud . A computer you rent from Amazon by the hour. NOVA runs on exactly one.
EBS	Elastic Block Store . A hard disk you attach to that computer. It survives the computer being restarted or replaced — this is where all your state lives.
Terraform	A tool that builds cloud infrastructure from a written description. You run one command and it creates the server, disk, permissions and logging. It can also show you what it is <i>about</i> to do before it does it, and undo it all later.
IAM	Identity and Access Management . AWS's permission system. An IAM role is a set of permissions something can wear. NOVA's role is deliberately tiny.
Permissions boundary	A ceiling on an IAM role. Even if someone later attaches a policy granting more, the boundary still refuses. A seatbelt for permissions.
SSM	Systems Manager (Session Manager) . Gets you a command line on the server <i>without</i> opening any port to the internet. This is why NOVA's server has no inbound access at all.
Bedrock	Amazon's service for calling AI models. You are billed per token of text in and out.
S3	Simple Storage Service . Amazon's file storage. Optionally where a customer's document library lives, mirrored into NOVA's knowledge base.
CloudWatch	Where AWS keeps logs and metrics. NOVA's container logs go here.
Secrets Manager	Encrypted storage for passwords and API keys. NOVA may read secrets under one named prefix and nowhere else.
KMS	Key Management Service . Holds the encryption key for the disk and the logs. The customer can supply their own — and for a professional-services firm, should.

Term	What it actually is
IMDSv2	A hardened way for the server to read its own identity. Required here. It closes a well-known attack where a bug in a web app is used to steal the server's AWS credentials.
Image digest	A fingerprint of an exact image, like sha256:fd57.... A tag can be moved to point at different content; a digest cannot. Always deploy by digest.
Tenant bundle	The folder of YAML files describing one customer: their agents, policy, knowledge sources, channels and automations. NOVA's input.
Apply	The command that reads the tenant bundle and writes the runtime configuration. It overwrites derived files, which is why edits belong in the bundle and never in the generated output.

17.3 The 36 steps

Before you start

Steps 1–17 build infrastructure. Steps 18–29 configure NOVA. Steps 30–36 prove it works. **Step 30 — running a real AI task — has never been performed in any deployment.** Budget real time for it and expect to find problems; that is what first deployments are for.

01 Create an AWS account

What you're doing. Sign up at aws.amazon.com. Use the **customer's** account, not yours.

Why. In the model you are selling, the customer owns the infrastructure and the data. It is also the answer to their biggest objection.

You should see. A root login and a 12-digit account number.

What can go wrong. Using your own account 'to make it easier' — you become the data processor and the whole security story collapses.

How to fix it. Start again in the customer's account. Do not migrate later; it is worse than starting over.

Security. Turn on MFA for the root user immediately, then never use root again. Create an admin IAM user for the work.

02 Choose an AWS region

What you're doing. Pick one region and use it for everything. For a UK customer: **eu-west-2 (London)**.

Why. Data residency. A UK accountancy firm's clients' records should sit in the UK, and a partner will ask. Region also affects which Bedrock models are available and what they cost.

You should see. A region code you will type into Terraform and the AWS console.

What can go wrong. Resources scattered across regions — invisible costs and a data-residency answer you cannot give.

How to fix it. Delete and rebuild in the right region. Terraform makes this cheap, which is one of the reasons to use it.

Security. Check Bedrock model availability in eu-west-2 *before* promising a specific model. Availability differs by region.

03 Install the local tools

What you're doing. On your own machine: Docker, the AWS CLI v2, Terraform, and the AWS Session Manager plugin.

Why. These four do everything: build the package, talk to AWS, build the infrastructure, and get a shell on the server.

You should see. Each responds to `--version`.

What can go wrong. Terraform v1 syntax against an older binary; Session Manager plugin forgotten until step 18.

How to fix it. Install the plugin now, not when you need it.

Security. Configure the CLI with an IAM user that has admin in this account, not root. Never commit credentials to a repository.

04 Build the NOVA Docker image

What you're doing. `deploy/docker/build.sh`

Why. The Terraform module pulls an image; it does not build one. This is the package.

You should see. A tag of the form `nova-control-plane:0.1.0-g<commit>` and a sha256 digest. About 304 MB.

What can go wrong. A dirty working tree, which appends `-dirty.<hash>` to the tag. That is the script working correctly — but do not deploy it.

How to fix it. Commit your work and rebuild. `build.sh --print-tag` tells you the tag for the tree you are on.

Security. The script also writes `:local`. That tag is mutable and must never be deployed.

05 Test the image locally

What you're doing. `deploy/docker/validate-local.sh <tag>`

Why. 49 checks: two tenants side by side, authentication, RBAC, isolation, CSP, graceful shutdown, and that no secret is in the image. Ten minutes here saves a day later.

You should see. 49 checks, 49 passed, 0 failed.

What can go wrong. Anything other than 49/49.

How to fix it. Stop. Do not push. A failure here is a failure in the customer's account too, just harder to see.

Security. Includes eight checks specifically for secrets in the image. Never skip them because you are in a hurry.

06 Create the ECR repository

What you're doing. In the customer's account, in your chosen region, create a private repository named `nova-control-plane`.

Why. Somewhere private to keep the image. The Terraform module does not create it — it only grants permission to pull from it.

You should see. A repository URI like `<account>.dkr.ecr.eu-west-2.amazonaws.com/nova-control-plane`.

What can go wrong. Creating it as public. Creating it in a different region from the server.

How to fix it. Delete and recreate. Private, same region.

Security. Private. Turn on **tag immutability** so a tag can never be silently repointed at different content. Turn on scan-on-push.

07 Authenticate Docker to ECR

What you're doing. `aws ecr get-login-password --region eu-west-2 | docker login --username AWS --password-stdin <account>.dkr.ecr.eu-west-2.amazonaws.com`

Why. ECR is private, so Docker needs a token to upload.

You should see. Login Succeeded.

What can go wrong. Token expired (they last 12 hours). Wrong region. Wrong AWS profile.

How to fix it. Re-run it. If it still fails, check `aws sts get-caller-identity` — you are probably in the wrong account.

Security. Never paste the password into a command line; the pipe above keeps it out of your shell history.

08 Push the image

What you're doing. Tag the local image with the ECR URI and `docker push` it.

Why. Gets the package into the customer's account so the server can download it.

You should see. Layers uploading, then a digest printed at the end.

What can go wrong. Pushing the `:local` tag. Pushing a `-dirty` tag.

How to fix it. Push the immutable tag from step 4. Delete anything wrong from ECR.

Security. This is the first time NOVA code enters the customer's account. Everything after this runs from what you pushed here.

09 Get the immutable image digest

What you're doing. From the push output, or: `aws ecr describe-images --repository-name nova-control-plane --image-ids imageTag=<tag>`

Why. A tag is a label that can be moved. A digest is the content itself. Deploying by digest means the server can only ever run the exact bytes you tested.

You should see. `sha256:` followed by 64 hex characters.

What can go wrong. Using the tag in `image_uri` instead. It will work, and then one day it will run something you did not test.

How to fix it. Use `<repo>@sha256:<digest>`.

Security. This is a supply-chain control. It is the difference between a reproducible deployment and a hopeful one.

10 Configure the tenant bundle

What you're doing. Copy `nova/examples/acme` and edit it for this customer: organization, identity, agents, policy, knowledge, channels.

Why. This is the customer's workforce, written down. It is also the thing the Control Center edits — nothing here is a one-off script.

You should see. `python -m nova validate <bundle>` passes.

What can go wrong. A permission not defined in the policy; a knowledge source id typo; an agent delegating to an agent that does not exist. All caught at load, by design.

How to fix it. Read the error — it names the file and the field.

Security. Never put a credential in the bundle. Six secret-shaped keys are refused at parse time at any nesting depth, and that refusal is deliberate.

11 Configure principals

What you're doing. `nova token new <name> --role admin` for each person. Write the printed lines into `control-principals.yaml` on the state volume.

Why. Who may use the Control Center, and at what level. Admins change things; viewers watch.

You should see. A token printed **once**, and a YAML entry containing only its SHA-256 digest.

What can go wrong. Losing the token. It is genuinely unrecoverable — NOVA never stores it.

How to fix it. Mint a new one and replace the entry.

Security. The command does not write the file — you do. Without a principals file any loopback caller is a local admin, and a non-loopback bind is refused outright. Give staff **viewer**; partners get admin.

12 Review the Terraform variables

What you're doing. Copy `terraform.tfvars.example` and fill in: `tenant_id`, `region`, `vpc_id`, `subnet_id`, `image_uri`, and optionally `bedrock_model_ids`, `secret_prefix` and `integrations`.

Why. These are the only decisions the module needs. Everything else has a safe default.

You should see. A `tfvars` file with no placeholder left in it.

What can go wrong. Leaving AWS's documented example values (111122223333, `vpc-0123...`) in place.

How to fix it. Replace every one. Terraform will not catch this — they are syntactically valid.

Security. The subnet needs **outbound internet** (a NAT gateway, or VPC endpoints). Do not solve that by putting the server in a public subnet with a public IP.

13 terraform init

What you're doing. `terraform init` in `deploy/aws`.

Why. Downloads the AWS provider and sets up state tracking.

You should see. Terraform has been successfully initialized!

What can go wrong. No internet; a provider version conflict.

How to fix it. `terraform init -upgrade`.

Security. For anything beyond a first test, configure **remote state** in S3 with locking. Local state on a laptop is a single point of failure for the customer's infrastructure.

14 terraform validate

What you're doing. `terraform validate`

Why. Checks the configuration is internally consistent before it talks to AWS.

You should see. Success! The configuration is valid.

What can go wrong. A typo in a variable name.

How to fix it. Fix what it names and re-run.

Security. Free and instant. Always run it.

15 terraform plan

What you're doing. `terraform plan -out=tfplan`

Why. Shows exactly what will be created, changed or destroyed. Nothing happens yet.

You should see. A list ending in `Plan: 13 to add, 0 to change, 0 to destroy`.

What can go wrong. Credentials for the wrong account. A subnet that is not in the stated VPC.

How to fix it. `aws sts get-caller-identity` to confirm which account you are pointing at.

Security. Never run **apply** without a saved plan. Applying an unreviewed plan in a customer's account is the fastest way to lose their trust.

16 Read the plan properly

What you're doing. Actually read it. All 13 resources.

Why. This is the last moment before anything exists. In a customer's account, it is also the artefact you can show them.

You should see. 1 EC2 instance · 1 EBS volume + attachment · 1 KMS key + alias · 1 CloudWatch log group · 1 security group with egress 443 only and no ingress · 2 IAM roles · 2 role policies · 1 permissions boundary policy · 1 attachment · 1 instance profile.

What can go wrong. Any ingress rule. A larger instance than agreed. A missing KMS key.

How to fix it. Do not apply. Fix the variables.

Security. Confirm there is no ingress rule and that IMDSv2 is required. Those two lines are most of the host's security posture. Save the plan output as deployment evidence.

17 terraform apply

What you're doing. terraform apply tfplan

Why. Builds it.

You should see. Two to four minutes, then outputs including the instance id.

What can go wrong. Service quota exceeded; a region that does not offer the instance type; insufficient IAM permissions on your own user.

How to fix it. Read the error — AWS names the quota or the permission. Quota increases are a support request and can take a day; do this before a customer is watching.

Security. You have now created billable resources in the customer's account. Set up the billing alarm in §18 **today**, not next week.

18 Connect using SSM

What you're doing. aws ssm start-session --target <instance-id> --region eu-west-2

Why. Gets you a shell on the server without any open port.

You should see. A shell prompt.

What can go wrong. TargetNotConnected — usually the instance has not finished booting, or the subnet has no route out to reach the SSM service.

How to fix it. Wait two minutes. Then check the subnet's outbound route or its VPC endpoints.

Security. This is the only way in, and that is the design. No SSH key exists. No bastion. Every session is logged by AWS. Do not 'temporarily' add an SSH rule.

19 Verify the EC2 host

What you're doing. Check the volume is mounted: df -h /var/lib/nova

Why. State must be on the encrypted EBS volume, not the instance's own disk.

You should see. An XFS filesystem mounted at /var/lib/nova.

What can go wrong. Not mounted — user_data failed part-way.

How to fix it. sudo cat /var/log/cloud-init-output.log and read the failure.

Security. If state is on the instance disk instead, everything is lost when the instance is replaced. Check this properly.

20 Verify Docker and the container

What you're doing. sudo docker ps

Why. Confirms the image pulled and the container is running.

You should see. One container from your ECR digest, status healthy after about 20 seconds.

What can go wrong. ImagePullBackOff or an authentication error — the instance role could not pull from ECR. A restart loop.

How to fix it. Check the ECR repository policy and that image_uri names the right account. For a restart loop: sudo docker logs <id> — most often 'no tenant bundle'.

Security. The image ships no tenant configuration and refuses to serve without one. Terraform does not place the bundle; you do. Decide how — synced from S3 at boot, or written during provisioning.

21 Verify NOVA's health

What you're doing. sudo docker exec <id> /usr/local/bin/nova-healthcheck, or curl the health route from inside the container.

Why. Confirms the control plane is accepting connections.

You should see. Exit code 0. HTTP 200, 401 or 403 all count as healthy.

What can go wrong. Expecting only 200 and concluding it is broken.

How to fix it. Nothing to fix. **401 is healthy** — with a TLS proxy declared, the control plane stops trusting loopback, so an unauthenticated probe correctly gets 401. That proves both the server and its auth layer are working.

Security. This is liveness, not readiness. It answers before a tenant bundle is valid — deliberately, so a typo in a policy file does not cause a restart loop.

22 Reach the Control Center securely

What you're doing. Port-forward over SSM: `aws ssm start-session --target <id> --document-name AWS-StartPortForwardingSession --parameters '{"portNumber":["8787"],"localPortNumber":["8787"]}'` then open `http://localhost:8787`.

Why. There is no inbound access and the container publishes no port by default. This tunnels through SSM instead of opening anything.

You should see. The Control Center in your browser, on your own machine.

What can go wrong. **This step needs a decision that has not been made.** As shipped, the systemd unit publishes no ports, so nothing on the host can reach 8787 at all.

How to fix it. Either add `--publish 127.0.0.1:8787:8787` to the unit and set `NOVA_BIND_HOST=0.0.0.0` plus `NOVA_BEHIND_TLS_PROXY=1` and a principals file — or skip the dashboard and use `docker exec` for CLI work.

Security. **Never solve this by adding a security-group ingress rule.** A non-loopback bind is refused without a principals file AND a TLS statement — both refusals confirmed firing. That guard is protecting the customer.

23 Create the tenant

What you're doing. Place the bundle at `/var/lib/nova/bundle` and run `python -m nova apply`.

Why. Turns the declarations into runtime profiles on disk.

You should see. Profiles created, an audit log written, and honest warnings about anything not yet possible.

What can go wrong. Warnings that each agent 'cannot run yet' because credentials are missing.

How to fix it. That is correct behaviour, not an error. Step 25 fixes it.

Security. Apply **overwrites** derived files, including each agent's persona. Edit the bundle, never the generated output.

24 Create the agents

What you're doing. Through the Control Center, or by adding YAML to the bundle and applying.

Why. The workforce itself.

You should see. Agents visible on the Agents screen, each with a materialised profile.

What can go wrong. A permission the policy does not define; an id that is not a valid identifier.

How to fix it. The error names the field. Fix the bundle.

Security. Every agent write is admin-only and audited intent → committed.

25 Configure the model provider

What you're doing. Set the model in the agent's spec, and put the credential in `<profile>/ .env`.

Why. Without a credential an agent is configured and cannot run. NOVA reports this honestly rather than failing silently.

You should see. `nova doctor` stops saying 'cannot run yet'.

What can go wrong. Exporting the key into the host environment instead of the profile's own file.

How to fix it. Put it in the profile's `.env`.

Security. **This matters more than it looks.** A credential in the host environment is readable by every agent on that host — verified. Isolation holds only when credentials live in the per-agent store.

26 Configure knowledge

What you're doing. Declare corpora in `knowledge.yaml`, put documents in the directory (or upload through the Control Center), and reindex.

Why. What the agents can actually answer from.

You should see. Documents listed, an index built, chunk counts shown.

What can go wrong. '0 indexed' — usually the include globs do not match the files.

How to fix it. Check the patterns. The Control Center shows what the corpus accepts.

Security. The walk refuses symlinks and anything resolving outside the root. That is the security boundary of the whole capability — do not work around it.

27 Configure S3-backed knowledge (optional)

What you're doing. Add an `origin` block naming the bucket, prefix and region, then Sync.

Why. Lets the firm keep documents where they already are instead of copying them onto the host by hand.

You should see. A sync report: downloaded, unchanged, skipped-and-why.

What can go wrong. An access-denied error — the instance role has no S3 permission by default.

How to fix it. Grant `s3:ListBucket` and `s3:GetObject` on that one prefix, via the integrations variable. Nothing broader.

Security. A mirrored corpus is read-only from NOVA. Uploads to it are refused, because the next sync would delete them. The bucket is where documents are managed.

28 Configure channels (optional)

What you're doing. Declare the connection, grant it to named agents, and put the credential in the agent's `.env`.

Why. How agents send and receive messages.

You should see. The channel listed with its granted agents.

What can go wrong. Assuming a platform works because it is in the catalogue.

How to fix it. Connect it and send a real message before telling a customer it is available.

Security. Zero of the 22 platforms are field-validated. Also: a webhook platform needs a publicly reachable HTTPS endpoint, and this deployment has no inbound access by design. That is a real decision, not a checkbox.

29 Create a governed automation

What you're doing. Declare it in the bundle or create it through the Control Center.

Why. Recurring work — the Records Chaser's Monday chase.

You should see. The automation listed, with its next run time.

What can go wrong. It is listed, the next run time passes, and **nothing happens**.

How to fix it. That is the known gap. The cron ticker lives inside the runtime's gateway; there is no standalone daemon. A gateway process must be running.

Security. The Automations screen leads with scheduler liveness for exactly this reason. Prove a scheduled job actually fires before any customer depends on one.

30 Run a real AI task

What you're doing. Submit an objective, or send a message on a connected channel, and watch a worker pick it up.

Why. The single most important step in this list. Until this happens, everything else is configuration.

You should see. A task on the Work board moving through states, a real model response, tokens recorded on the Usage screen.

What can go wrong. No provider configured; the policy denying a tool the agent needs; the model id not available in this region.

How to fix it. Read the Activity screen — every refusal is logged with the rule that caused it.

Security. This has never been done in a real deployment. Nothing in NOVA has called a real model provider. Budget real time for this step; expect to find things.

31 Verify logs

What you're doing. CloudWatch Logs in the console, and the Activity screen in the Control Center.

Why. Two different logs answering two different questions: the container's output, and what the governance layer decided.

You should see. Container logs arriving in CloudWatch; agent and error logs readable per agent in the Control Center.

What can go wrong. Nothing in CloudWatch — the log driver or the IAM permission.

How to fix it. Check the unit's `--log-driver awslogs` configuration and the role policy.

Security. Set a log retention period now. Logs default to never expiring and become a surprisingly large line on the bill.

32 Verify the audit trail

What you're doing. `/var/lib/nova/home/nova/audit.jsonl`

Why. The record of who did what. This is the artefact that makes the governance claim true.

You should see. Paired records: `intent` then `committed` or `failed`, each with the authenticated human and a correlation id.

What can go wrong. The actor showing as a service name rather than a person — that means someone is sharing a token.

How to fix it. One principal per human. That is the entire point.

Security. **Show this file to the customer during the pilot.** It is the most persuasive artefact you have, and it is worth more than any slide.

33 Test restart and recovery

What you're doing. Restart the container. Then stop and start the EC2 instance.

Why. Proves state survives, and that you can recover from a host problem.

You should see. The container comes back healthy; agents, automations, provenance and audit history all intact.

What can go wrong. State lost — it was not on the mounted volume.

How to fix it. Go back to step 19.

Security. Also run `nova backup` and **practise a restore**. A backup you have never restored is not a backup. Note that backups exclude credentials by default, deliberately.

34 Test tenant isolation

What you're doing. Confirm this deployment holds exactly one tenant and its state is under one root.

Why. One tenant per deployment is the isolation guarantee. There is no multi-tenancy.

You should see. One bundle, one tenant id stamped on every audit event.

What can go wrong. Attempting to serve a second customer from the same host.

How to fix it. Do not. Deploy a second stack. It is one Terraform apply.

Security. **Never co-locate two customers.** The architecture does not support it and the repository says so plainly.

35 Test permissions

What you're doing. Create a viewer principal. Confirm it can read the agents list and is refused on policy, decisions, budget and every write. Then attempt a denied tool call and watch the refusal appear in Activity.

Why. Proves RBAC and the policy layer are live in *this* deployment, not just in tests.

You should see. 403s where expected, and a logged refusal naming the rule.

What can go wrong. A viewer seeing something they should not.

How to fix it. Stop and investigate before go-live. Do not proceed.

Security. **Do this in front of the customer.** A demonstrated refusal is worth more than any security page on your website.

36 Record the deployment evidence

What you're doing. Save: the plan output, the image digest, the Terraform outputs, the 49-check result, screenshots of steps 30–35, and the AWS account and region.

Why. This becomes the repeatable runbook for customer two, and the evidence pack for the first security questionnaire you are asked to complete.

You should see. A dated folder per customer.

What can go wrong. Not doing it, then rebuilding the knowledge from memory for the next customer.

How to fix it. Do it while it is fresh — the same day.

Security. Store it where the customer's credentials are not. Evidence packs get emailed.

18 AWS cost control

Structure, not prices

Why there are no prices in this section

AWS pricing changes, varies by region, and is only authoritative on AWS's own pages — which were **unreachable from the environment this document was built in** (the network egress policy blocks `aws.amazon.com`). Quoting figures from memory or from a third-party blog would be exactly the kind of fabricated number this document exists to avoid. What follows is the cost *structure*, which does not change, plus the official URLs to price it against on the day you quote.

18.1 What generates a bill

Service	What you pay for	How big it is, relatively	How to control it
EC2	The server, per hour it exists — running or not, unless stopped	Largest fixed cost. A general-purpose instance running continuously.	Right-size it. Start small; NOVA's control plane is not compute-heavy. Consider a Savings Plan once the size is settled.
EBS	The disk, per GB per month, whether used or not	Small but permanent — and it survives the instance.	Provision what you need. Snapshots also cost; set a retention policy.
ECR	Image storage per GB/month, plus data transfer	Very small. The image is ~304 MB.	A lifecycle policy to expire untagged images. Otherwise every build accumulates.
CloudWatch Logs	Ingestion per GB, then storage per GB/month	The classic surprise. Logs default to never expiring.	Set a retention period on day one. 30 or 90 days. This is the single highest-value cost action in this table.
S3	Storage, requests, and retrieval	Small for a document library; the request count can surprise you.	Lifecycle rules. Note that NOVA's sync compares ETags and does not re-download unchanged objects, which keeps request volume down.
Bedrock	Per 1,000 tokens in and out, per model	The only variable that can grow without limit. Depends entirely on agent activity and model choice.	See §18.3. This is the one to watch weekly.
Data transfer	Mostly outbound to the internet; NAT gateway processing	Modest here — no public traffic — but a NAT gateway has its own hourly charge plus per-GB.	VPC endpoints for ECR, S3, SSM and CloudWatch instead of a NAT gateway can be cheaper. Price both.
KMS	Per key per month, plus API requests	Small and fixed.	One key. A customer-managed key is worth the small cost for this buyer.
Secrets Manager	Per secret per month, plus API calls	Small.	Few secrets. Do not put one secret per agent per variable in here unless you need to.

18.2 Price it here, on the day

Official AWS pricing pages — the only authoritative source

EC2 on-demand: <https://aws.amazon.com/ec2/pricing/on-demand/>

EBS: <https://aws.amazon.com/ebs/pricing/>

ECR: <https://aws.amazon.com/ecr/pricing/>

CloudWatch: <https://aws.amazon.com/cloudwatch/pricing/>

S3: <https://aws.amazon.com/s3/pricing/>

Bedrock: <https://aws.amazon.com/bedrock/pricing/>

KMS: <https://aws.amazon.com/kms/pricing/>

Secrets Manager: <https://aws.amazon.com/secrets-manager/pricing/>

Data transfer: https://aws.amazon.com/ec2/pricing/on-demand/#Data_Transfer

Build the estimate in the **AWS Pricing Calculator** (<https://calculator.aws>), save it, and send the customer the link. It is a shareable URL, it is authoritative, and it moves the cost conversation from your word to Amazon's.

“It runs on AWS” does not mean it is free — say so early

A non-technical buyer often hears “your own AWS account” as “no extra cost”. Correct that in the first conversation, not in month two when the bill arrives. The honest framing: **“You pay Amazon directly for the server and the AI usage. That's typically a few hundred pounds a month at this size, it's in your account so you can see every line of it, and I'll help you set a budget alarm on day one.”** Then actually help them set it.

18.3 Beginner checklist for avoiding a surprise bill

- | | |
|---|--|
| ☐ Billing alerts on, with an email that someone reads | ☐ AWS Budgets: a monthly budget with an 80% forecast alert |
| ☐ CloudWatch log retention set (30 or 90 days) — not 'never expire' | ☐ ECR lifecycle policy to expire untagged images |
| ☐ S3 lifecycle rules if a document library is large | ☐ Instance right-sized — start small, measure, then resize |
| ☐ Nothing left running from testing (old instances, orphaned volumes, old snapshots) | ☐ Orphaned EBS volumes deleted — they bill after the instance is gone |
| ☐ NAT gateway vs VPC endpoints priced both ways before choosing | ☐ Bedrock usage reviewed weekly for the first month |
| ☐ Cost allocation tags on every resource, so the customer can see NOVA's share | ☐ A named person on the customer's side who sees the bill |

There is no hard spending ceiling in NOVA — plan around it

This is stated in §3.8 and it bears repeating here because it is a money question. Run-budget seconds trigger a soft wrap-up message; nothing terminates. Token and cost figures on the Usage screen are **observed, not enforced**. The controls that actually exist are: AWS Budgets alarms, the per-run tool-call ceiling in the compiled policy, a conservative model choice, and looking at the bill weekly. **Never tell a customer NOVA caps their spend.**

19

Production checklist

Two gates, before and during

19.1 Before any customer

These are about *you*, not about a deployment. Do not take a first customer's money with more than two of these unticked.

- | | |
|--|--|
| ☐ Security review of the deployment completed and written down | ☐ Tenant isolation verified on real infrastructure |
| ☐ IAM policies reviewed line by line by someone other than the author | ☐ TLS decision made and documented (proxy, or certificate) |
| ☐ Authentication proven: principals file, hashed tokens, no shared logins | ☐ RBAC proven: viewer refused on every admin route |
| ☐ Secrets: none in the image, none in the bundle, none in the audit log | ☐ Backup taken AND a restore practised on a fresh host |
| ☐ Logging: CloudWatch delivery confirmed, retention set | ☐ Monitoring: health check alarm, disk-space alarm, billing alarm |
| ☐ Cost controls: budget, alerts, log retention, ECR lifecycle | ☐ Recovery: instance replacement tested with state intact |
| ☐ Upgrade procedure written — new digest, re-apply, verify | ☐ Rollback procedure written and tested — previous digest, redeploy |
| ☐ Deployment runbook written from a real deployment, not from this document | ☐ Capability schedule (§3) current, and matching every sales claim you make |

19.2 Customer deployment

- ☐ Discovery: workflow observed, baseline measurement started
- ☐ AWS account access confirmed, region agreed (eu-west-2 for UK)
- ☐ ECR repository created, private, tag-immutable
- ☐ Terraform plan reviewed WITH the customer and saved as evidence
- ☐ Tenant bundle placed and validated
- ☐ Model provider configured; credentials in the profile .env, not the host
- ☐ Integrations: S3 origin and/or channels, each granted per agent
- ☐ Policy compiled; a denied action demonstrated to the customer
- ☐ Automations declared; scheduler liveness confirmed firing
- ☐ Validation: health, logs, audit trail, restart, isolation, permissions
- ☐ Monitoring and budget alarms live before go-live, not after
- ☐ Day-60 ROI review booked in the calendar now
- ☐ Architecture agreed with the customer's IT person or MSP
- ☐ IAM: admin user created, root MFA on, root not used again
- ☐ Image pushed; deployment pinned to the digest, not a tag
- ☐ Terraform applied; 13 resources confirmed; no ingress rule
- ☐ Principals file written; partners admin, staff viewer
- ☐ Knowledge corpora loaded and indexed; counts verified
- ☐ Agents created; personas written with the customer
- ☐ Green/Yellow/Red agreed, signed, and reflected in the policy
- ☐ A real AI task executed end to end and watched
- ☐ Training delivered — especially the approvals queue habit
- ☐ Deployment evidence pack saved and dated

20

12-month commercial roadmap

Ordered by customer value

Nothing here is a feature somebody imagined. Every product item either closes a gap identified in §3 and §23, or was asked for by a real customer — and in months 1–3 there are no customers, so everything is gap-closing. That is the correct order.

MONTH 1 ● Field validation — nothing else matters

Deploy to a real AWS account. Push to ECR. Prove ECR pull, KMS volume, awslogs, SSM, IMDSv2. Execute one real model call through Bedrock. Prove one scheduled job fires. Settle the Control Center exposure decision.

MONTH 2 ● Make it repeatable, and start the list

Second deployment from the runbook with no improvisation. Backup and restore drill. Build the 100-prospect list. Record the 90-second demo. Governance page live.

MONTH 3 ● First pilot

Outreach at volume. 8–12 discovery calls. First paid pilot signed and started with a two-week baseline. First security questionnaire answered — keep the answers.

MONTH 4 ● Deliver, and measure honestly

Pilot weeks 2–6. Weekly reviews. Sample and grade 30 outputs a week. Fix what the first real workload breaks — it will break something.

MONTH 5 ● Convert and prove

Day-60 ROI review. Convert to standard pricing. Write the case study with their numbers and their quote. Ask for two referrals in that meeting.

MONTH 6 ● Second customer

Sell with the case study. Raise price to standard tier. Onboard the sales partner (§15) with the capability schedule attached.

MONTH 7 ● Customer success becomes a function

Monthly reviews for both customers. First upsell — a second workflow. Task detail and run history surfaced, because customers will ask.

MONTH 8 ● Security posture

Formal security questionnaire pack. DPIA template for customers. Penetration test. Start the SOC 2 conversation if mid-market is the target.

MONTH 9 ● Third and fourth customers

Repeatable sale. Partner produces their first independently-sourced deal. Vertical-specific content compounding.

MONTH 10 ● Vertical specialisation

A pre-built accountancy tenant bundle: standard agents, standard policy, standard Green/Yellow/Red. Deployment time falls from weeks to days.

MONTH 11 ● Channel

Two or three referral partners: practice-software consultants, accountancy-focused MSPs, fractional FDs. Their relationship, your product.

MONTH 12 ● Assess honestly

Five or six customers, a repeatable deployment, a repeatable sale, one vertical understood. Decide: hire a salesperson (§15.2), raise, or stay deliberately small.

20.1 By function

Function	Months 1–3	Months 4–6	Months 7–9	Months 10–12
Product	Close the field-validation gaps. No new features.	Fix what the first real workload exposes.	Task detail, run history, comments — customers ask for these.	Accountancy bundle template. Semantic retrieval only if customers actually hit the BM25 limit.
Engineering	SQLite 3.53 build stage. Exposure decision. Gateway process for the scheduler.	Repeatable deployment. Backup/restore drill. Upgrade and rollback.	Observability. Alerting. On-call basics.	Hardening for mid-market. Multi-deployment management tooling.
Sales	100-prospect list. Demo. First 8–12 calls.	First pilot. First conversion. Case study.	Partner onboarded. Customers 3 and 4.	Repeatable process documented. Decide on the first hire.
Marketing	Governance page. LinkedIn founder cadence. Demo video.	Case study. ROI method published.	Vertical content. Conference talk or webinar.	Retargeting test. First paid experiment, small.
Customer success	n/a	Weekly pilot reviews. Output sampling.	Monthly reviews. First upsell. Referral asks.	Health scoring. Renewal process. QBRs.
Security	Internal review. IAM audit.	First questionnaire. DPIA template.	Penetration test. Incident response plan.	SOC 2 readiness if mid-market is the target.
Hiring	Nobody.	Nobody.	Commission-only sales partner.	Assess against \$15.2. Probably a delivery engineer before a salesperson.

21

Revenue scenarios

Illustrative arithmetic, not forecasts

Illustrative scenario, not a financial forecast
NOVA has **no customers and no revenue**. Every figure in this section is arithmetic on a stated assumption. Its purpose is to show the shape of the business and where the margin sits — not to predict anything. Do not put these in an investor deck as projections.


21.1 The assumptions, stated once

Assumption	Value used	Basis
Blended implementation fee	£6,000	Between the founding (£4,500) and standard (£7,500) tiers — \$7
Blended monthly fee	£1,200	Between founding (£950) and standard (£1,450)
Your infrastructure cost per customer	£0	The customer pays AWS directly, in their own account. This is the BYOC model and it is a real structural advantage
Your delivery cost per customer per month	£150	Assumed: support time, monitoring, your own tooling amortised
Churn	0%	Assumed away for simplicity. This is unrealistic and the biggest weakness in these numbers
Implementation gross margin	~50%	Assumed: the rest is your delivery time at a notional rate

21.2 Scenarios

Customers	Setup revenue (one-off)	MRR	ARR	Your cost/mo	Gross margin on recurring
5	£30,000	£6,000	£72,000	£750	£5,250/mo (87.5%)
10	£60,000	£12,000	£144,000	£1,500	£10,500/mo (87.5%)
25	£150,000	£30,000	£360,000	£3,750	£26,250/mo (87.5%)
50	£300,000	£60,000	£720,000	£7,500	£52,500/mo (87.5%)
100	£600,000	£120,000	£1,440,000	£15,000	£105,000/mo (87.5%)

ILLUSTRATIVE SCENARIO — not a forecast. Assumes £6,000 setup, £1,200/month, £150/month delivery cost per customer, zero churn, and that the customer pays AWS directly.



ILLUSTRATIVE SCENARIO — not a forecast. One stated assumption: £1,400/month per customer and no churn, with customers landing on the ramp shown. Reaching 13 customers in 12 months would be a strong first year and is not a prediction.

21.3 What these numbers hide

The number looks like	The reality
87.5% gross margin on recurring	True only if delivery stays at £150/customer/month. The first three customers will cost far more than that in your own time, and that time is the real constraint on growth.
Setup revenue scales linearly	Each implementation is real work. Without the month-10 vertical bundle, customer 25 costs roughly what customer 5 did — and you only have so many weeks.
Zero churn	The single most unrealistic assumption here. Model 10–20% annual churn for an SMB product with no switching cost and see what the curve does.
Customers pay AWS, so infrastructure is free to you	True, and a genuine structural advantage — but it also means a customer’s model bill is a churn risk you do not control and cannot cap (\$18.3).
100 customers is a business	100 single-tenant deployments is 100 things to upgrade, monitor and support. Without the month-12 deployment tooling, that is an operations problem long before it is a revenue achievement.

22

Competitive positioning

Where NOVA actually sits



Two axes a buyer in this market actually chooses on. Positions are my assessment of each category, not a scored benchmark of named products.

22.1 Category by category

Category	Examples	What they do better	Where NOVA wins
Workflow automation	Zapier, Make, Power Automate	Vast connector libraries, instant setup, trivially cheap, enormous communities.	They move data between apps on fixed rules. They do not reason over documents, do not have a policy layer, do not have an approvals queue, and run in someone else's cloud. Different product for a different job — say so rather than competing.
AI copilots	Microsoft 365 Copilot, Salesforce Agentforce, Gemini for Workspace	Already in the customer's stack, per-seat pricing, no deployment, huge trust advantage from the brand.	Confined to their own vendor's data and surfaces. Assist a person rather than hold a workflow. This is your most common real competitor — the answer is "Copilot helps your people work faster; NOVA does the job."
Hosted AI agent platforms	Various 2026 agent-building platforms	Fast to start, growing fast, good developer experience, no infrastructure work.	The customer's data goes to their cloud. Governance is usually developer-facing rather than business-facing. For a firm holding client confidential records, "in your own account" is the whole conversation.
AI automation agencies	The many UK agencies in £7.1	Cheap to start, flexible, sell on outcomes, move fast.	They build bespoke workflows on someone else's platform, usually with no governance layer and no audit trail. When the agency goes quiet, the customer has an unmaintainable thing. This is who you are competing with on price , and governance is the differentiator.
Enterprise AI platforms	watsonx Orchestrate, UiPath, Appian, and similar	Real governance, self-hosted options, compliance certifications, established sales machinery.	Priced and scoped for a 5,000-person organisation. A 40-person practice cannot buy them, cannot implement them, and is not worth their sales cycle. That gap is your market.

On Palantir comparisons — don't

It will be tempting, because the shape rhymes: governed, deployed into the customer's environment, audit-heavy. **NOVA is not comparable to Palantir** and claiming it is will end a credible conversation instantly. Different scale, different problem, different maturity, thousands of engineer-years apart. If someone raises it, the honest answer is: "Same instinct about governance, completely different scale. We do one thing for firms of 20 to 150 people." That answer earns respect; the alternative does not.

22.2 What is actually differentiated, and how defensible it is

Differentiator	How real	How defensible
Runs in the customer's own AWS account	Real — the Terraform module exists and is tested. Unproven on real AWS.	Strong. Most competitors are structurally hosted and cannot follow without rebuilding their business model.
Governed AI workforce — policy, approvals, per-agent permissions	Real and enforced. Fail-closed policy, compiled per agent, checked at every tool call.	Strong for 12–24 months. It is genuinely hard to retrofit, and the hard-won lessons (a plugin that was installed and never consulted) are not obvious from outside.
Write-ahead audit trail	Real and live-proven locally. Intent before, outcome after, authenticated human.	Moderate. Not hard to build; it is hard to remember to build, and harder to retrofit into a product that did not start with it.
Configurable workforce, not code	Real — agents, personas, permissions and schedules are all editable in the browser and all write to the bundle.	Moderate. Everyone is building this. Your version's advantage is that edits are durable through apply, which is less obvious than it sounds.
Enterprise Control Center	Real. 13 screens, strict CSP, RBAC, tenant-scoped API.	Weak alone. UI is copyable. It matters as the delivery vehicle for the governance, not on its own.
Built on an open runtime	Real and honest.	Double-edged. You inherit a large capable runtime for free — and a dependency you do not control. See §23.

23

Risks and honest limitations

Everything that could stop this

This section is mandatory reading before any sales conversation. Nothing is hidden and nothing is softened. The last column is the one that matters: whether it blocks the first customer.

23.1 Technical and validation risks

Limitation	Severity	Impact	Action	Blocks #1?
No AWS field validation. ECR pull, instance profile, KMS volume mount, awslogs delivery, SSM access and ExternalId integration roles are configured and have never run.	CRITICAL	First deployment becomes debugging in front of a paying customer. Any of the six could fail.	Deploy to your own AWS account this month. \$24 action 1.	YES
No real provider validation. No agent has called Bedrock or any hosted model. The live worker run used a scripted server.	CRITICAL	Token accounting, cost estimation, streaming, tool-call accumulation and error handling are all unproven against a real provider.	Execute one real Bedrock call end to end. \$24 action 2.	YES
Scheduled execution never observed firing. The cron ticker lives inside the runtime's gateway; there is no standalone daemon.	CRITICAL	A deployment can hold a perfectly correct schedule that nothing runs. The Records Chaser depends on this entirely.	Run a gateway process and watch a job fire. Until then, scope pilots to the Librarian and Client Desk. \$24 action 3.	YES
Control Center exposure undecided. No ingress, no published port, SSM gives a host shell — which cannot reach port 8787.	CRITICAL	As shipped, the customer cannot open the dashboard at all.	Choose: loopback publish + SSM port-forward (recommended), or CLI-only. Document it. \$24 action 4.	YES
Zero channels field-validated. 5 read in source, 17 manifest-only.	HIGH	Any channel promised to a customer might not work. Email is the one the first offer depends on.	Field-validate email first, then one chat platform. Never put a logo on a slide before connecting it.	Partly
SQLite 3.40.1 in the image carries the upstream WAL-reset bug. The runtime detects it and degrades to journal_mode=DELETE.	MEDIUM	No corruption risk — the degradation is safe. What is lost is reader/writer concurrency, which matters for a busy 24/7 board.	Port the existing SQLite 3.53 builder stage from the root Dockerfile. Blocked in this environment by egress policy, not by difficulty.	No
Single-node architecture. State is SQLite on one attached volume.	MEDIUM	No horizontal scale, no HA. An instance failure is downtime until it is replaced.	Correct for this market. Be explicit in the SLA; do not promise HA.	No
No multi-tenancy. One tenant per deployment.	MEDIUM	Every customer is a separate stack to deploy, upgrade and monitor. This is an operations cost that grows linearly.	Accept it — it is also the security story. Build multi-deployment tooling around month 12, not multi-tenancy.	No
Backup/restore never drilled on real infrastructure.	HIGH	A backup you have never restored is not a backup.	Drill it on a fresh host in month 2, before customer one goes live.	Partly
Thin observability. Health check is liveness only; no alerting beyond what CloudWatch gives you.	MEDIUM	You will find out a customer's deployment is broken when they tell you.	Health alarm, disk alarm, billing alarm on day one of the first deployment.	No

Limitation	Severity	Impact	Action	Blocks #1?
Positive tool scoping not enforced. toolsets/allow are recorded, not enforced by the adapter.	MEDIUM	A narrower grant than the customer believes. Denials ARE enforced, so the practical exposure is bounded.	Use tools.deny for anything that matters. Never claim allow-listing as an enforced control.	No
BM25 keyword retrieval, no embeddings.	LOW	Poor recall where the customer's question uses different vocabulary from their documents.	Sell it accurately. Revisit only if a real customer hits the limit.	No

23.2 Commercial and operational risks

Limitation	Severity	Impact	Action	Blocks #1?
No hard cost ceiling. Structurally unavailable — run-budget seconds are a soft wrap-up, not a kill.	HIGH	A runaway loop on an expensive model produces a bill nobody capped. Also a renewal risk you do not control.	Customer pays AWS directly; AWS Budgets alarms; conservative model choice; per-run tool-call ceiling; weekly review. Never claim a cap.	No
Model costs are variable and the customer sees them.	MEDIUM	Bill shock in month two kills the relationship faster than a product fault.	Set expectations in the first conversation. Set the alarm on day one. Review weekly for the first month.	No
Integration complexity. Every firm's practice software is configured differently.	MEDIUM	Implementation overruns eat the fee and your calendar.	Scope one workflow. Quote integrations separately. Say no to bespoke in the first three customers.	No
Customer security reviews. No SOC 2, no ISO 27001, no penetration test, no production history.	HIGH	Blocks mid-market and kills regulated sectors outright.	Target SMB first, where the partner is the reviewer. Build the questionnaire pack from the first one you answer.	No
UK GDPR and ICO expectations. The ICO's January 2026 agentic-AI report flags purposes set too broadly and unfettered access to data and systems; most agentic deployments touching personal data warrant a DPIA.	HIGH	Client records are personal data. A firm that has not done a DPIA is exposed, and so are you.	Turn this into a sales asset. Provide a DPIA template. Narrow purposes and per-agent grants are literally what the ICO asks for — you are selling the control, not the risk.	No
Professional-body caution. ICAEW is publicly flagging over-reliance on AI as a reputational threat to the profession.	MEDIUM	A partner may decline on principle, or wait for guidance.	Lead with governance and approvals. Keep the first workflows away from anything resembling advice.	No
Support burden. One person, N single-tenant deployments, no on-call.	HIGH	Caps how many customers you can hold. Roughly 5–8 before it breaks.	Price for it. Business-hours SLA only. Invest in tooling at month 12, not more customers.	No
Sales cycle. Professional-services firms are seasonal — January is a write-off, and so is the run-up to any filing deadline.	MEDIUM	Three to four months from first contact to signature, with dead months in it.	Build the pipeline before you need it. Avoid launching outreach into January.	No

Limitation	Severity	Impact	Action	Blocks #1?
Upstream dependency on Hermes. ~600k lines you do not control.	HIGH	A breaking upstream change, an abandoned project, or a licence change is an existential event. A capability already vanished once: the profile allowlist key NOVA compiled into was removed upstream and had to be probed for at runtime.	Pin the version. Keep the patch budget near zero. Maintain the compatibility manifest. Watch upstream releases. Contribute so you have standing.	No
Key-person concentration. One person holds product, sales and delivery.	HIGH	Illness or burnout stops the company. It also caps growth before money does.	Document everything as you go. The runbook and the capability schedule are the first two artefacts that survive you.	No
No customers, no case study, no references.	HIGH	The hardest sale you will ever make is the first one.	Founding-customer pricing with the case study written into the contract. Be honest that they are first — some buyers value it.	Partly

The four that actually block the first customer

AWS field validation · real provider execution · scheduled execution · the Control Center exposure decision.

Everything else on these two pages can be managed, disclosed or priced around. These four cannot, and they are the first four actions in §24. They are also, between them, probably two to three weeks of work — which is the most encouraging sentence in this document.

24

What I should do next

Fifteen actions, in order

Ordered by what unblocks the next thing, not by what is most interesting. The first four are the blockers from §23 and nothing else should be started until they are done.

01 Deploy to a real AWS account — your own This week

Push the image to ECR, run terraform apply, and prove the six unexercised things: ECR pull, instance profile, KMS volume mount, awslogs delivery, SSM access, IMDSv2. Use your own account so the first failures are private. **Nothing else on this list matters until this is done.**

02 Execute one real model call through Bedrock Same week

Configure a Bedrock model, put the credential in a profile's .env, submit an objective, and watch a worker complete it. Check token accounting on the Usage screen against the AWS bill. This is the step that turns a control plane into a product.

03 Prove a scheduled job actually fires Week 2

Run a gateway process, create an automation, and watch it execute. Record the execution ledger row. Until you have seen this, do not sell the Records Chaser.

04 Settle and document Control Center access Week 2

Choose the loopback-publish plus SSM port-forward route, configure it in the systemd unit, write the principals file, and confirm the refusals fire when you get it wrong. Write it into the runbook.

05 Field-validate email as a channel Week 3

The first offer depends on email. Connect it, send and receive a real message, and promote it from source_read to field_validated with the evidence. Then one chat platform.

06 Drill backup and restore on a fresh host Week 3

Take a backup, destroy the instance, rebuild, restore, verify agents, automations, knowledge and audit history. Time it, and write the number into the SLA.

07 Write the deployment runbook from a real deployment Week 4

Not from §17 of this document — from what actually happened, including what went wrong. This is the artefact that makes customer two economic.

08 Confirm the vertical and build the 100-prospect list Week 4

UK accountancy practices, 20–150 staff. By hand, with a named partner and a specific pain hypothesis per row. The research is the qualification.

09 Build the demo and the 90-second video Week 5

A seeded demo tenant that looks like a real practice. Rehearse §13 until it runs to time without notes. Record the video from the real Control Center.

10 Publish the governance page and start the LinkedIn cadence Week 5

One page, before the rest of the website. Three posts a week from your personal account. The audience takes eight weeks to build, so start it before you need it.

11 Contact all 100 and run the discovery calls Weeks 6–8

Follow §9. Expect roughly 20 replies, 10 calls, 6 demos. Record every lost reason.

12 Sign and run the first paid pilot Weeks 9–10

Founding pricing. Two-week baseline first — refuse to start without it. Case study written into the agreement.

13 Measure, review at day 60, and convert Weeks 11–18

Their numbers. Three honest outcomes on the table. Ask for two referrals in that meeting, while the number is in front of them.

14 Write the case study and deploy customer two from the runbook Month 5–6

Named firm, their figures, their quote. Then prove the runbook works by using it without improvising. Raise to standard pricing.

15 Onboard the sales partner with the capability schedule Month 6+

Commission-only, structured per §15, with the one-page annex of what may and may not be claimed. Not before you have closed three yourself.

If you only do three things

Actions 1, 2 and 3. They are probably two weeks of work between them, and they move NOVA from “a well-built control plane that has never run anywhere real” to “a product with a deployment.” Every commercial sentence in this document gets easier to say honestly the moment they are done — and none of them can be said properly until they are.

25

Sources and references

Where every external number came from

Repository claims are cited inline by file and symbol throughout §1–§3 and are verifiable at commit b9319677. External figures are below. Where a source is self-published by a party with a commercial interest in the number, that is marked.

25.1 UK market and AI adoption

Claim used	Source
UK AI adoption ~35% (10+ employees), 29% all bands, ~1.6 technologies per adopter	Office for National Statistics, <i>Artificial intelligence in UK businesses: 2023 to 2026</i> , published 20 July 2026, ons.gov.uk/businessindustryandtrade/business/businessservices/articles/artificialintelligenceinukbusinesses/2023to2026
54% of UK SMEs use AI in some form; 11% use it extensively to automate operations; 95% report no workforce reduction	British Chambers of Commerce, <i>Half of SMEs Using AI</i> , March 2026
12% of AI-using businesses report increased revenue; 75% report productivity gains	ONS / DSIT, 2026
~5.68m UK SMEs (0–249 employees); 5.64m with 0–49 employees	UK Business Population Estimates 2025 (DBT), via House of Commons Library research briefing SN06152
~29,582 UK accounting & auditing businesses	IBISWorld, <i>Accounting & Auditing in the UK, 2025 (commercial research)</i>
~35,921 active UK accounting companies with recent filings	Firmbase industry list, 2026, derived from Companies House (<i>commercial aggregator — different definition from IBISWorld</i>)

25.2 UK accountancy sector

Claim used	Source
73% of UK accounting firms turning away clients for lack of staff; 74% believe workloads could push people out of the profession	Advancetrack, <i>Accounting Talent Index 2026</i>
95% of mid-tier firms expect increased AI use; 91% expect increased automation; 86% include AI in strategy; 17% believe they can assess AI's workforce impact; 66% investing in upskilling	ICAEW mid-tier firms research, 35 member firms with 11–249 principals, surveyed February–March 2026
ICAEW flagging misuse and over-reliance on AI as a reputational threat to the profession	ICAEW insights, 2026
Making Tax Digital for Income Tax live from 6 April 2026 for sole traders and landlords above £50,000 qualifying income	HMRC Making Tax Digital timetable, 2026
UK accounts assistant salary ~£23,700–£27,100 depending on source and geography	Multiple 2026 salary trackers; one cites ONS ASHE-derived data at ~£24,993. <i>Use the customer's own payroll figure in any ROI model.</i>

25.3 Pricing and commission benchmarks

Claim used	Source
UK single-workflow automation £1,000–£5,000 build, £200–£800/month support; connected multi-workflow £4,000–£12,000	UK AI automation agency pricing guides, 2026 (<i>self-published by agencies — directional only</i>)
SMB 2–3 workflows \$1,000–\$3,500/month; mid-market \$4,000–\$10,000/month; custom agent development from \$10,000	International AI automation agency pricing surveys, 2026 (<i>self-published</i>)
UK marketing agencies spending £350–£1,100 per seat/month on AI tooling	UK agency AI stack reporting, 2026 (<i>self-published</i>)
B2B sales commission 5–15% of deal value; SaaS new ARR 8–12%; renewals 2–8%; pure-commission independent representatives 15–30%	2026 B2B and SaaS commission benchmark surveys (Everstage, CaptivateIQ, Fullcast) (<i>vendor-published</i>)

25.4 Regulatory

Claim used	Source
ICO early views on agentic AI: risk of purposes set too broadly and unfettered access to data and systems; DPIA warranted where deployment is high-risk; accountability challenges extend to DPOs	Information Commissioner's Office report on agentic AI, January 2026. The ICO states it is early-stage thinking and not formal guidance — represent it that way to customers.
UK GDPR obligations (lawful basis, data minimisation, purpose limitation, accuracy, accountability) apply fully to AI systems	ICO guidance on AI and data protection

25.5 AWS

No AWS prices are quoted anywhere in this document

The official AWS pricing pages were unreachable from the build environment — the network egress proxy blocks `aws.amazon.com`. Rather than quote figures from memory or from a third-party blog, §18 gives cost structure and the official URLs. **Price every quote against those pages on the day, and build the estimate in the AWS Pricing Calculator so the customer gets Amazon's number rather than yours.**

Topic	Official source
EC2 on-demand pricing	<code>aws.amazon.com/ec2/pricing/on-demand/</code>
EBS pricing	<code>aws.amazon.com/ebs/pricing/</code>
ECR pricing	<code>aws.amazon.com/ecr/pricing/</code>
CloudWatch pricing	<code>aws.amazon.com/cloudwatch/pricing/</code>
S3 pricing	<code>aws.amazon.com/s3/pricing/</code>
Bedrock pricing	<code>aws.amazon.com/bedrock/pricing/</code>
KMS pricing	<code>aws.amazon.com/kms/pricing/</code>
Secrets Manager pricing	<code>aws.amazon.com/secrets-manager/pricing/</code>
Pricing calculator	<code>calculator.aws</code>
Session Manager (SSM) documentation	<code>docs.aws.amazon.com/systems-manager/latest/userguide/session-manager.html</code>
IMDSv2 documentation	<code>docs.aws.amazon.com/AWSEC2/latest/UserGuide/configuring-instance-metadata-service.html</code>
IAM permissions boundaries	<code>docs.aws.amazon.com/IAM/latest/UserGuide/access_policies_boundaries.html</code>

25.6 Repository — the primary source for every capability claim

Document	What it establishes
docs/audits/AWS_DEPLOYMENT_READINESS.md	Image, tag, ports, environment, volumes, AWS services, health check, the 49 local checks, and the ten remaining deployment requirements.
docs/audits/NOVA_HERMES_CAPABILITY_AUDIT.md	The ladder, and what the runtime provides versus what NOVA surfaces. Every row cites a file and symbol.
docs/audits/NOVA_CONTROL_CENTER_GAP_ANALYSIS.md	Gaps in order, with what would have to be true to call each validated.
docs/audits/PHASE_11_AUTOMATIONS.md, PHASE_12_GOVERNED_AUTOMATIONS.md	The scheduler, why create was withheld and then governed, and the ticker limitation.
docs/platform/PRODUCTION_READINESS_AUDIT.md	Twenty findings and their current status.
docs/platform/LIVE_RUN.md	The live worker run and the four seam defects only running it could find.
docs/PHASE_9.md, docs/NOVA_CHANNEL_SECURITY.md	The channel facade, the verification ladder, and the channel threat model.
CONTROL-CENTRE-EXTENSIONS-AUDIT.md	S3 knowledge mirroring, MCP and plugins — measured counts and the OAuth reality.
nova/capabilities/catalog.yaml	Machine-readable capability catalogue. The source of truth when this document and the repository disagree.
deploy/aws/, deploy/docker/	The Terraform module and the image build, including validate-local.sh.

When this document and the repository disagree, the repository is right. This playbook was generated against commit b9319677 on 15 September 2026. Regenerate it after any change that moves a capability up or down the ladder in §3 — and update the sales material and the partner capability schedule at the same time.